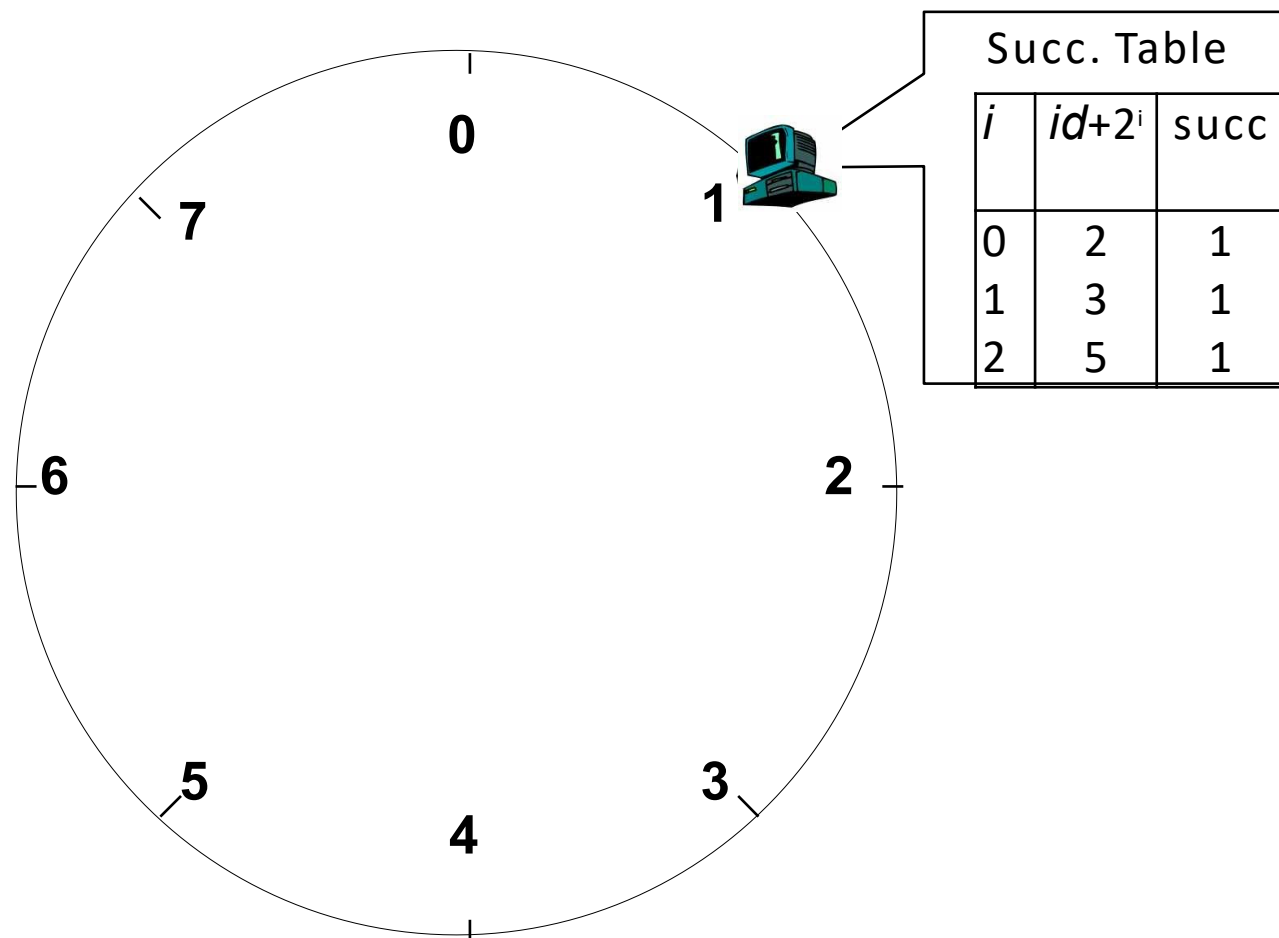


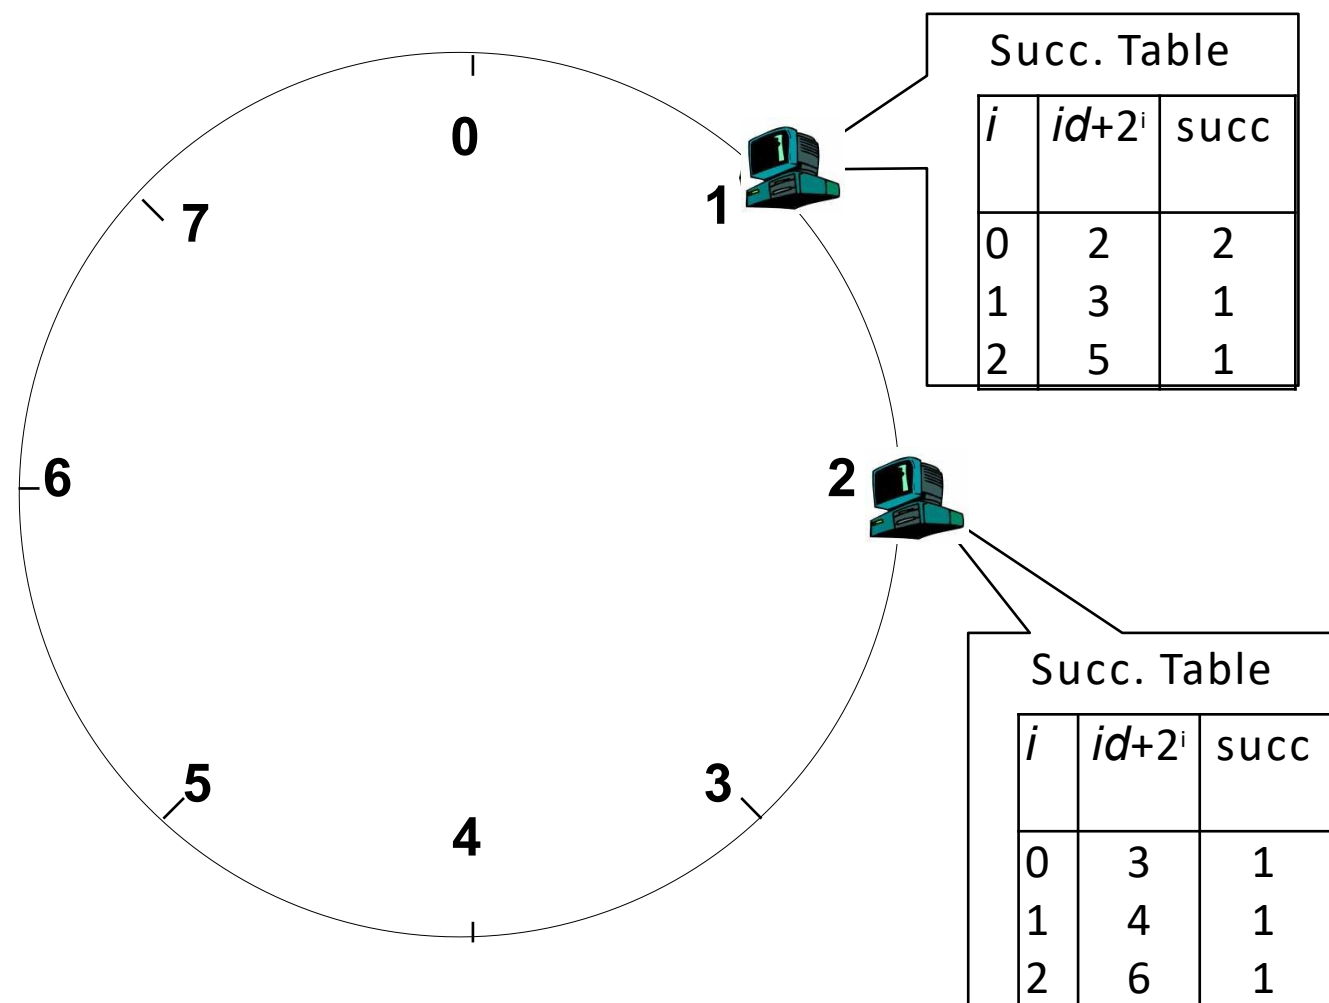
DHT: Chord Join

- Assume an identifier space $[0..8]$
- Node n1 joins



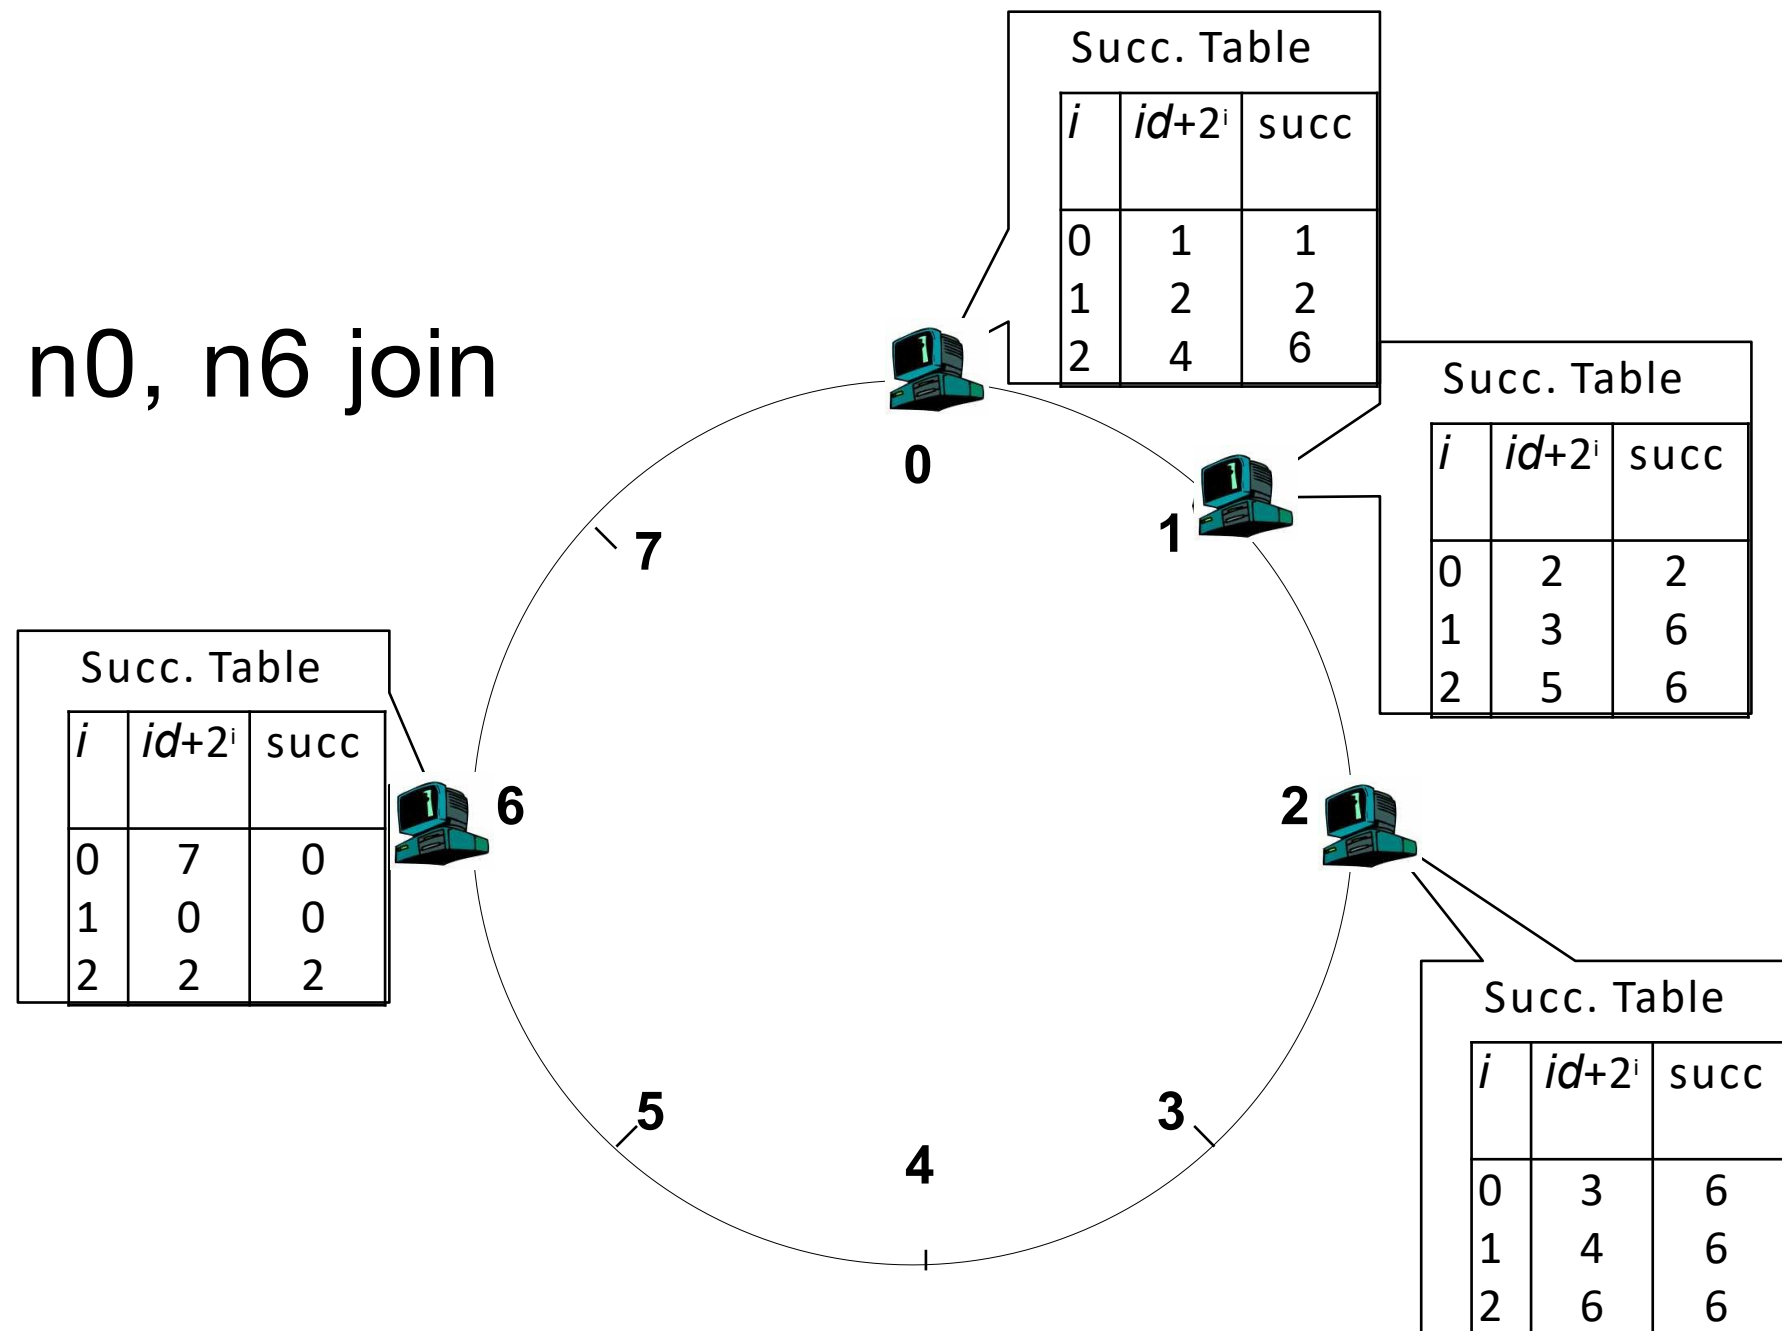
DHT: Chord Join

- Node n2 joins



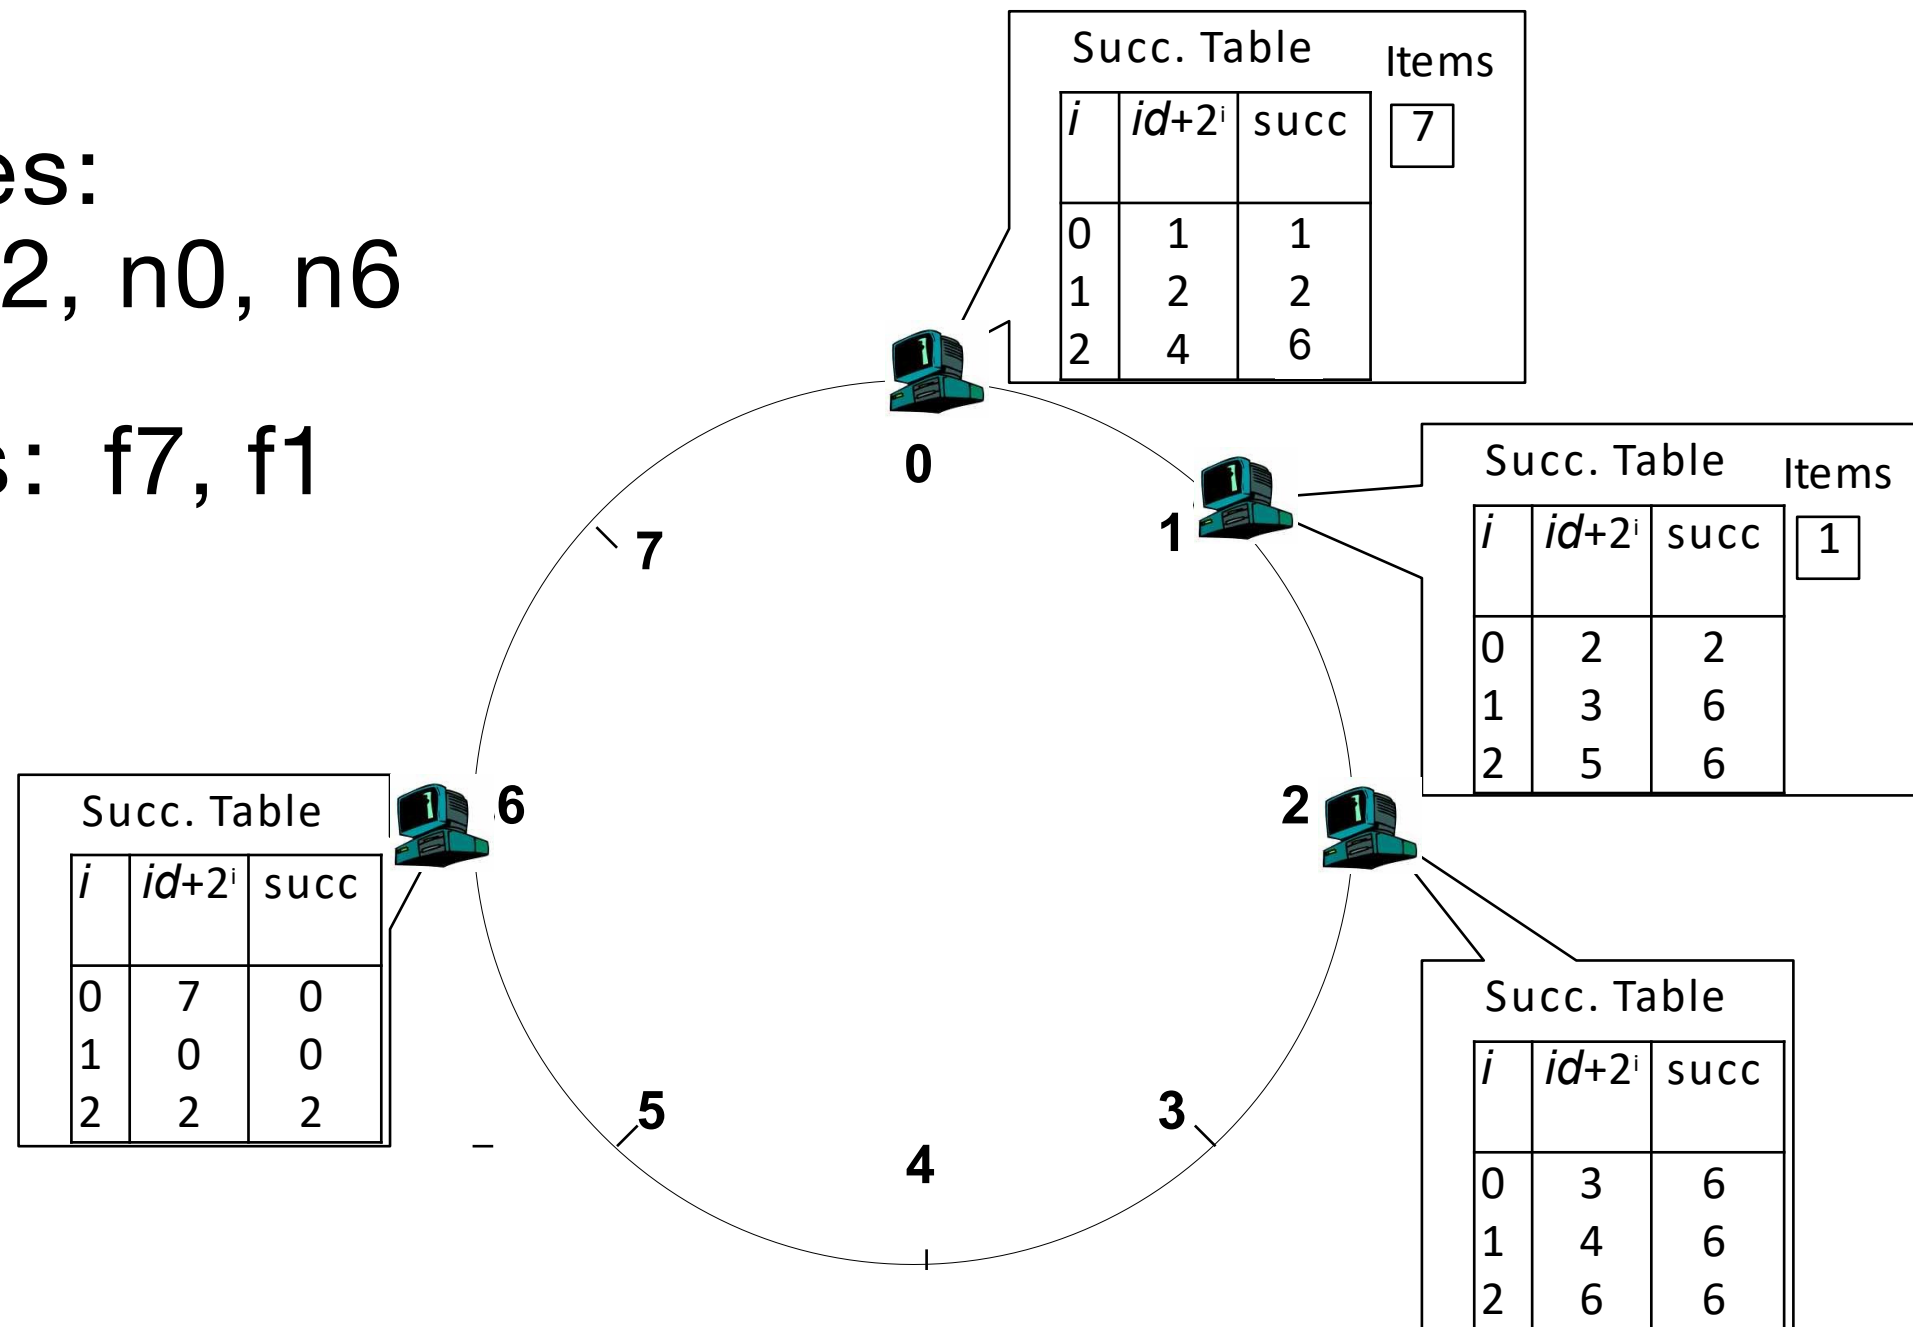
DHT: Chord Join

- Nodes n0, n6 join



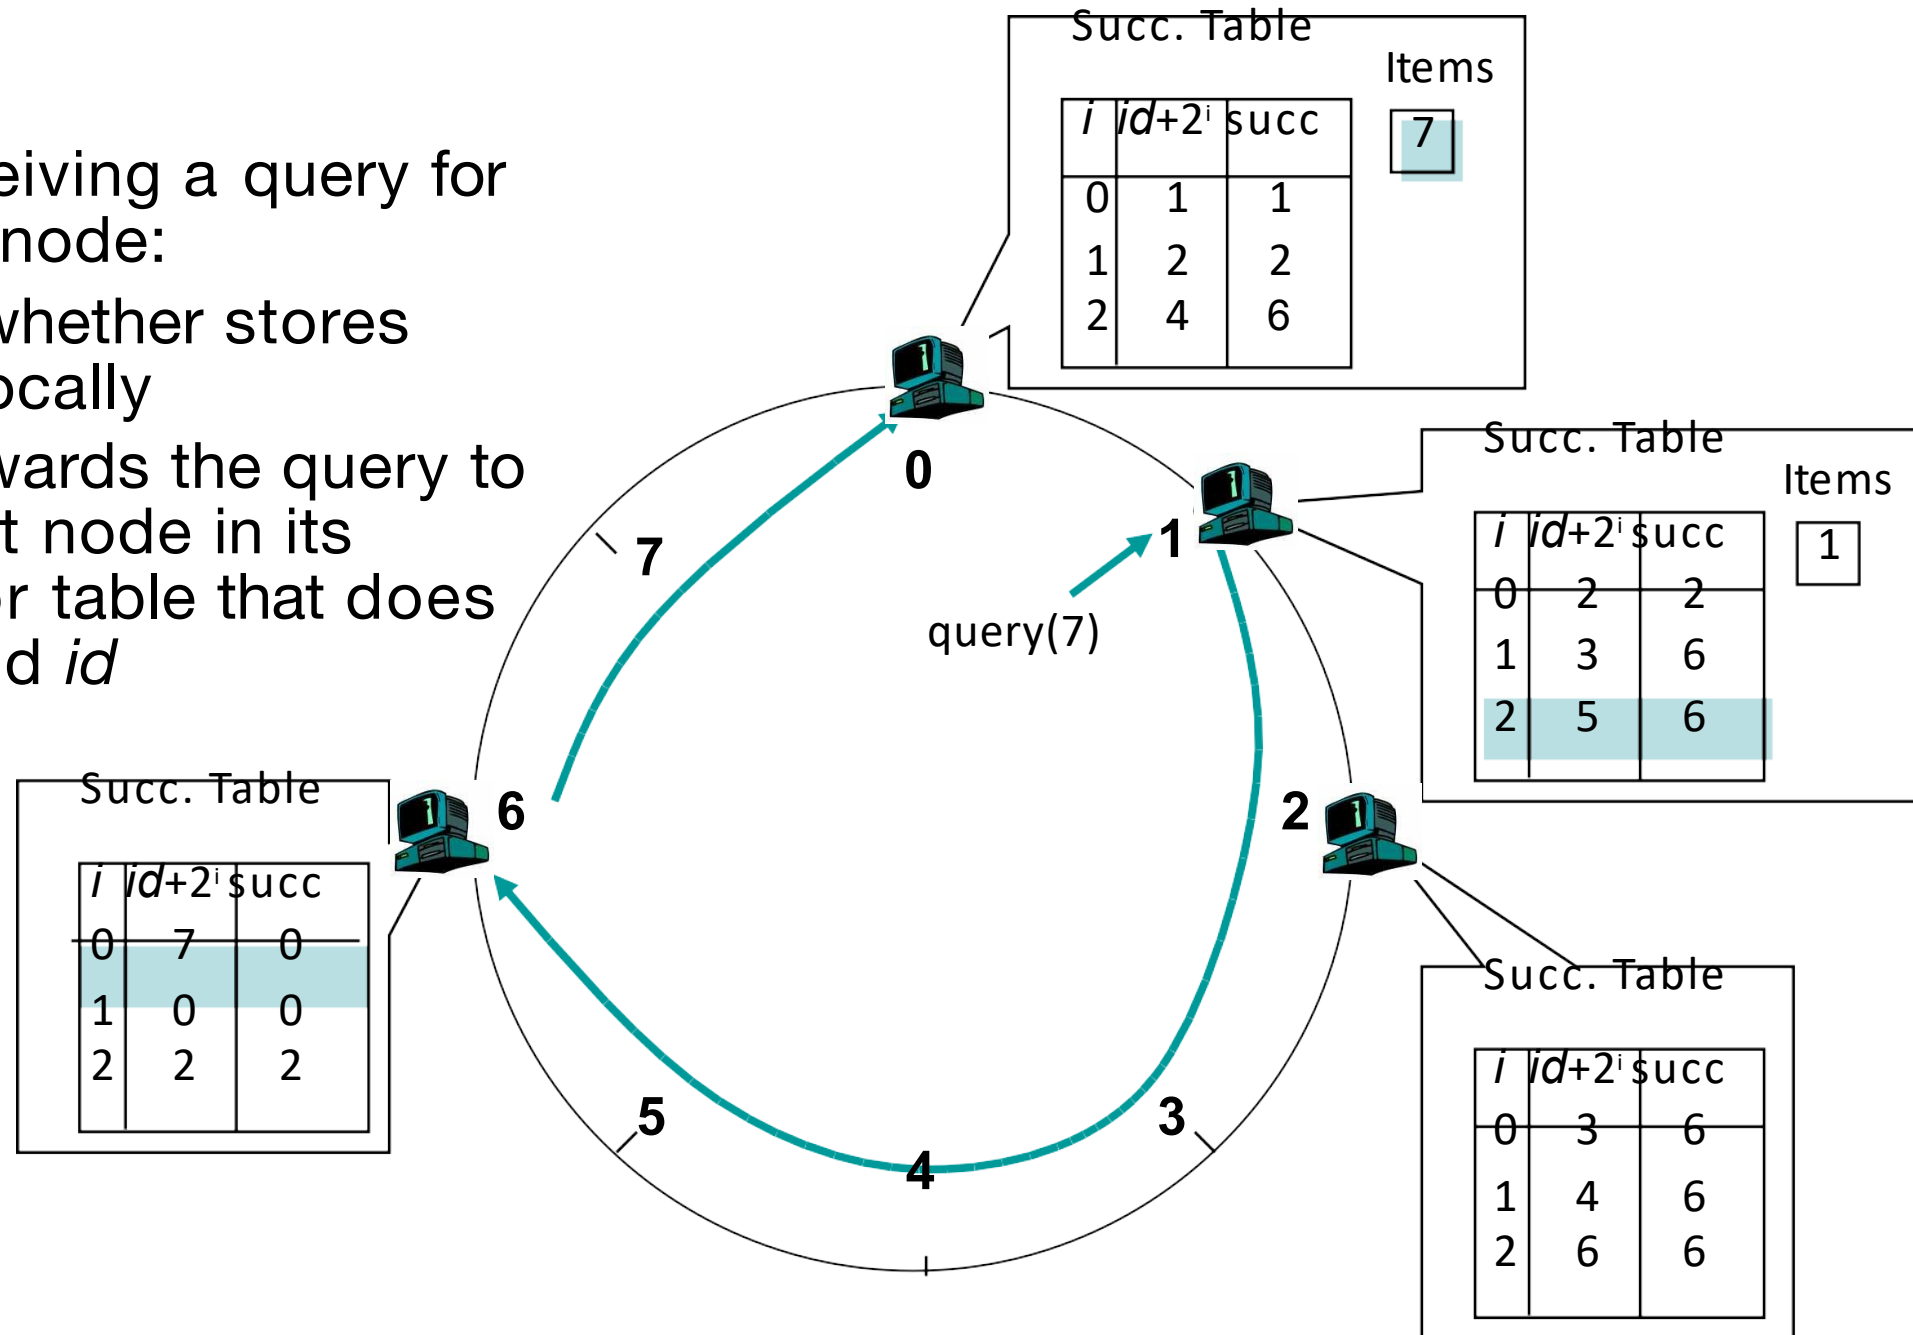
DHT: Chord Join

- Nodes:
n1, n2, n0, n6
- Items: f7, f1



DHT: Chord Routing

- Upon receiving a query for item id , a node:
- Checks whether stores the item locally
- If not, forwards the query to the largest node in its successor table that does not exceed id



DHT: Chord Summary

- ▶ Routing table size?
 - ▶ $\log N$ fingers
- ▶ Routing time?
 - ▶ Each hop expects to $1/2$ the distance to the desired id \Rightarrow expect $O(\log N)$ hops.

Case Study: Amazon Dynamo Key-Value Store

Amazon's workload (in 2007)

- **Tens of thousands** of servers in globally-distributed **data centers**
- **Peak load:** Tens of millions of customers
- **Tiered** service-oriented architecture
 - **Stateless** web page rendering servers
 - **Stateless** aggregator servers
 - **Stateful** data stores (e.g. **Dynamo**)
 - **put(), get():** values “usually less than 1 MB”

How does Amazon use Dynamo?

- **Shopping cart**
- **Session info**
 - Maybe “recently visited products” *etc.*?
- **Product list**
 - Mostly read-only, replication for high read throughput

Dynamo requirements

- **Highly available writes** despite failures
 - Despite disks failing, network routes flapping, “data centers destroyed by tornadoes”
 - **Non-requirement:** Security, viz. authentication, authorization (used in a non-hostile environment)
- **Low request-response latency.** focus on 99.9% SLA
- **Incrementally scalable** as servers grow to workload
 - Adding “nodes” should be seamless
- Comprehensible **conflict resolution**
 - High availability in above sense implies conflicts

Design questions

How is data **placed and replicated**?

How are **requests routed and handled** in a replicated system?

How to cope with temporary and permanent **node failures**?

Dynamo's system interface

Basic interface is a key-value store

- **get(k)** and **put(k, v)**
- Keys and values opaque to Dynamo

get(key) → value, **context**

- Returns one value or multiple conflicting values
- Context describes version(s) of value(s)

put(key, **context**, value) → “OK”

- **Context** indicates **which versions** this version supersedes or merges

Dynamo's techniques

Place replicated data on nodes with **consistent hashing**

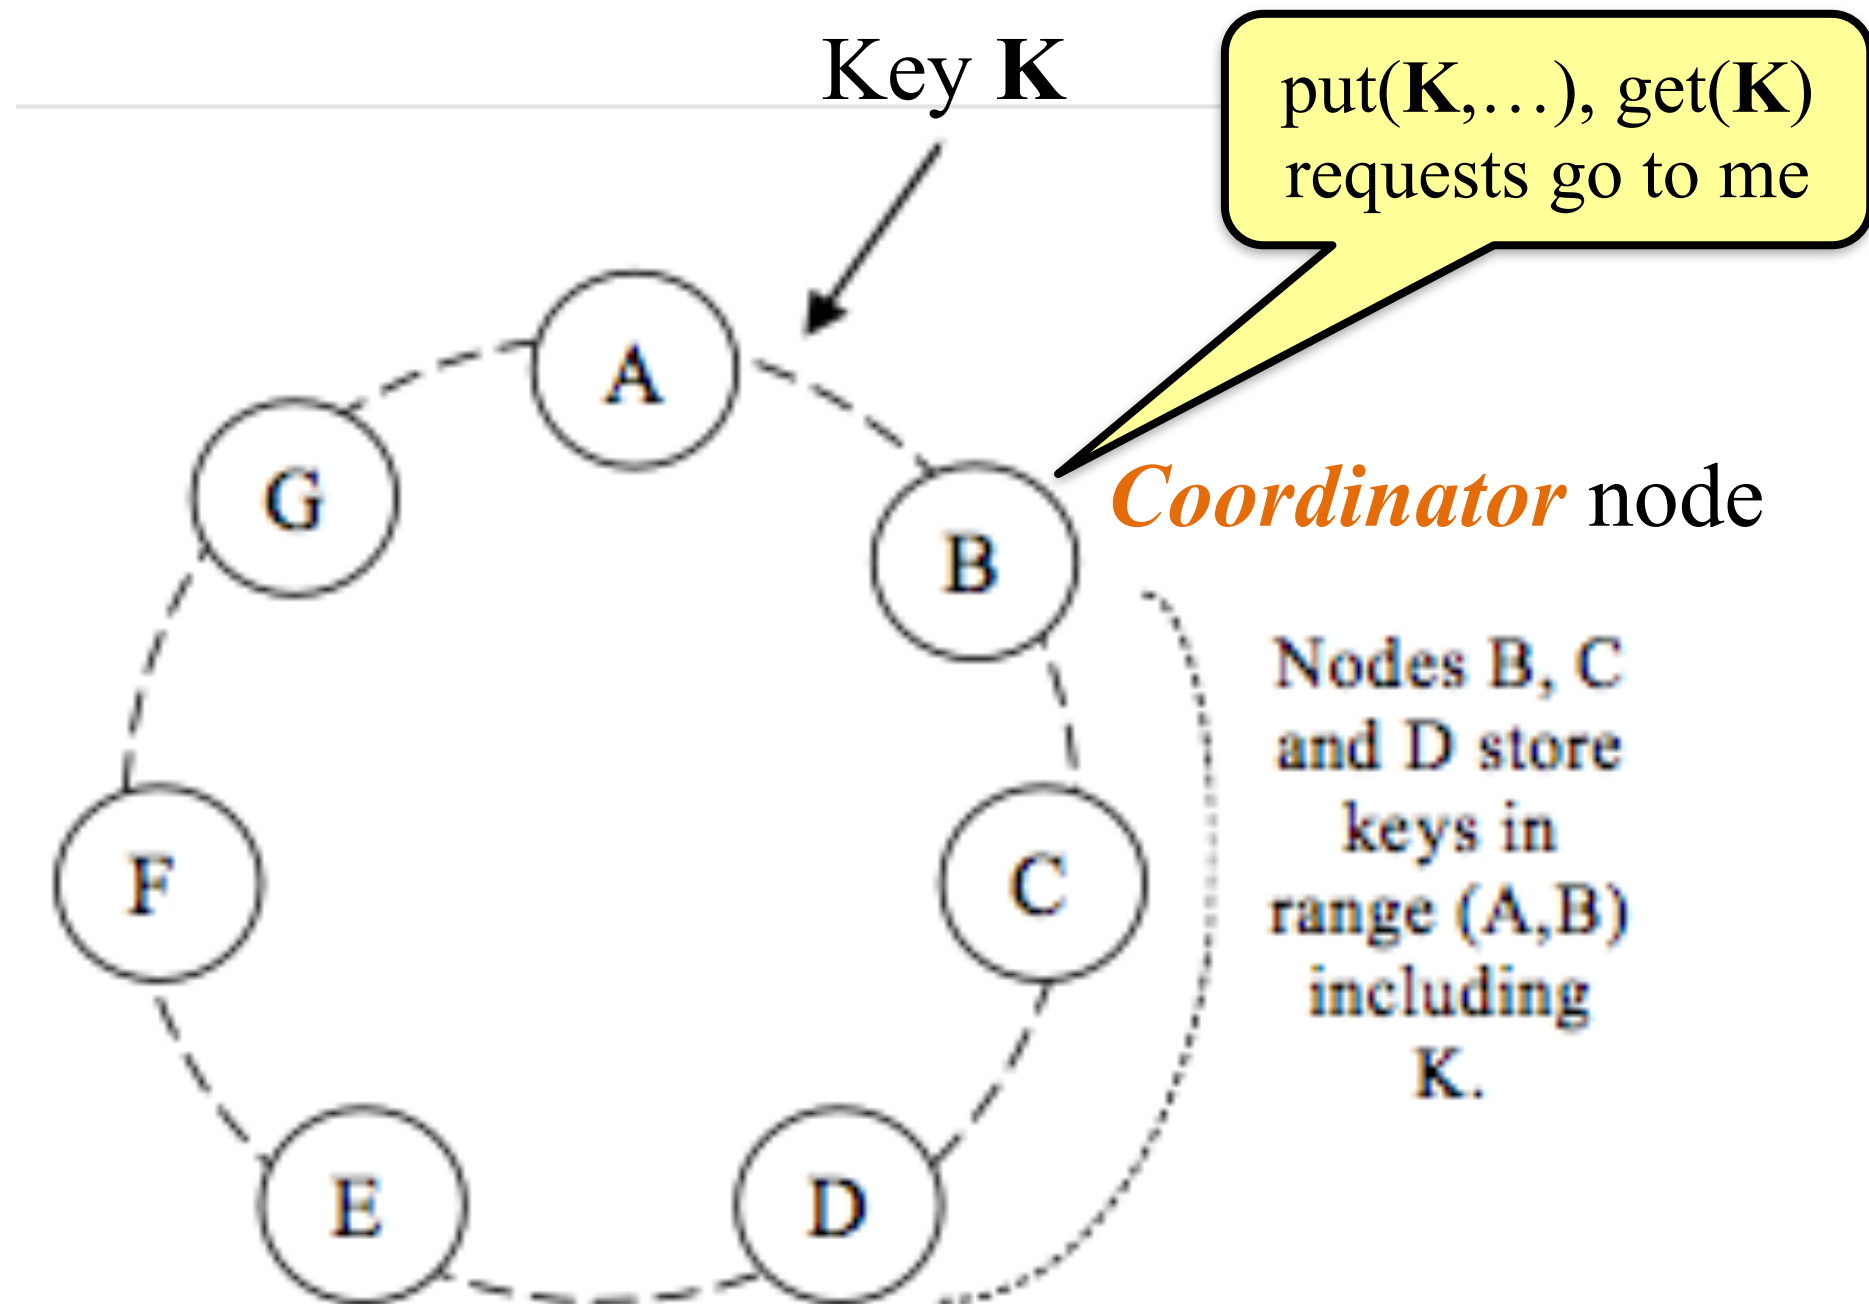
Maintain consistency of replicated data with **vector clocks**

- **Eventual consistency** for replicated data: prioritize success and low latency of writes over reads
 - » And availability over consistency (unlike DBs)

Efficiently **synchronize replicas** using **Merkle trees**

Key trade-offs: Response time vs. consistency vs. durability

Data placement



Each data item is **replicated** at N virtual nodes (e.g., $N = 3$)

Data replication

Much like in Chord: a key-value pair \rightarrow key's N successors (**preference list**)

- **Coordinator receives a put** for some key
- Coordinator then **replicates data onto nodes** in the key's **preference list**

Preference list **size** $> N$ to account for node failures

For robustness, the preference list **skips tokens** to **ensure distinct physical nodes**

Gossip and “lookup”

Gossip: Once per second, each node contacts a **randomly chosen other node**


- They **exchange their lists of known nodes**
(including virtual node IDs)

Each node **learns** which others handle all **key ranges**

- **Result: All nodes can send directly to any key's coordinator (“zero-hop DHT”)**
 - » **Reduces variability** in response times

Partitions force a choice between availability and consistency

Suppose three replicas are partitioned into two and one



If one replica fixed as master, no client in other partition can write

In Paxos-based primary-backup, no client in the partition of one can write

Traditional distributed databases emphasize consistency over availability when there are partitions

Alternative: Eventual consistency


Dynamo emphasizes **availability over consistency** when there are partitions

Tell client write complete when only some replicas have stored it

Propagate to other replicas in background

Allows writes in both partitions...but risks:

- Returning **stale data**
- **Write conflicts** when partition heals:



Mechanism: Sloppy quorums

If **no failure**, reap **consistency benefits** of single master

- Else **sacrifice consistency** to **allow progress**

Dynamo tries to store all values `put()` under a key on **first N live nodes** of coordinator's **preference list**

BUT to speed up `get()` and `put()`:

- Coordinator returns “success” for **put** when **$W < N$** replicas have completed **write**
- Coordinator returns “success” for **get** when **$R < N$** replicas have completed **read**

Sloppy quorums: Hinted handoff

Suppose coordinator **doesn't receive W replies** when replicating a put()

- Could return failure, but remember goal of **high availability for writes...**

Hinted handoff: Coordinator **tries further nodes** in preference list (**beyond first N**) if necessary

- Indicates the **intended replica node** to recipient
- **Recipient** will periodically try to forward to the **intended replica node**

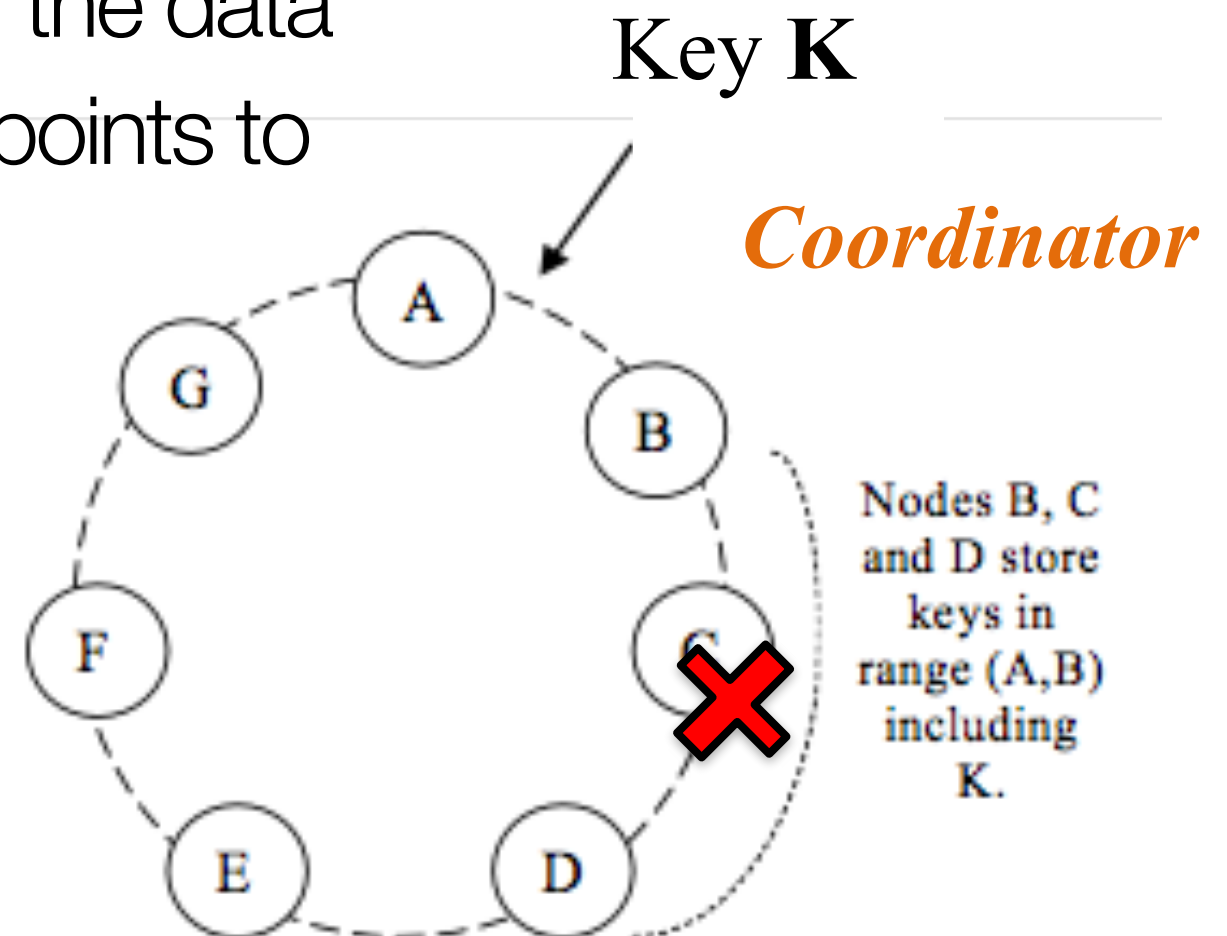
Hinted handoff: Example

Suppose **C fails**

- **Node E** is in **preference list**

- » Needs to receive replica of the data

- Hinted Handoff: replica at **E** points to node **C**



When **C comes back**

- **E** forwards the replicated data back to **C**

Sloppy quorums and get()s

Suppose coordinator **doesn't receive R replies** when processing a get()

- “ R is the min. number of nodes that must participate in a successful read operation.”
 - » Sounds like these get()s fail

Why not return whatever data was found, though?

- As we will see, consistency not guaranteed anyway...

Sloppy quorums and freshness

Common case given in paper: **$N = 3; R = W = 2$**

- With these values, **do sloppy quorums guarantee a `get()` sees all prior `put()`s?**

If no failures, **yes:**

- **Two writers** saw each `put()`
- **Two readers** responded to each `get()`
- Write and read **quorums must overlap!**

Sloppy quorums and freshness

Common case given in paper: **$N = 3, R = W = 2$**

- With these values, **do sloppy quorums guarantee a `get()` sees all prior `put()`s?**

With **node failures**, **no**:

- **Two nodes** in preference list **go down**
 - » `put()` replicated **outside preference list**
- **Two nodes** in preference list **come back up**
 - » `get()` occurs before they receive prior `put()`

Conflicts

Suppose **N = 3, W = R = 2**, nodes are named **A, B, C**

- 1st put(k, ...) completes on **A** and **B**
- 2nd put(k, ...) completes on **B** and **C**
- Now get(k) arrives, completes first at **A** and **C**

Conflicting results from **A** and **C**

- Each has seen a **different put(k, ...)**

Dynamo returns both results; what does client do now?

Conflicts vs. applications

Shopping cart:

- **Could take union** of two shopping carts
- What if second put() was result of user deleting item from cart stored in first put()?
 - » **Result: “resurrection” of deleted item**

Can we do better? Can Dynamo resolve cases when multiple values are found?

- **Sometimes.** If it can't, **application** must do so.

Version vectors (vector clocks)

Version vector: List of (coordinator node, counter) pairs

– e.g., [(A, 1), (B, 3), ...]

Dynamo stores a version vector with **each stored** key-value **pair**

Idea: track “ancestor-descendant” relationship between different versions of data stored under the same key **k**

Version vectors: Dynamo's mechanism


Rule: If vector clock comparison of $v1 < v2$, then the first is an ancestor of the second – **Dynamo can forget v1**

Each time a put() occurs, Dynamo increments the counter in the V.V. for the coordinator node

Each time a get() occurs, Dynamo returns the V.V. for the value(s) returned (in the “context”)

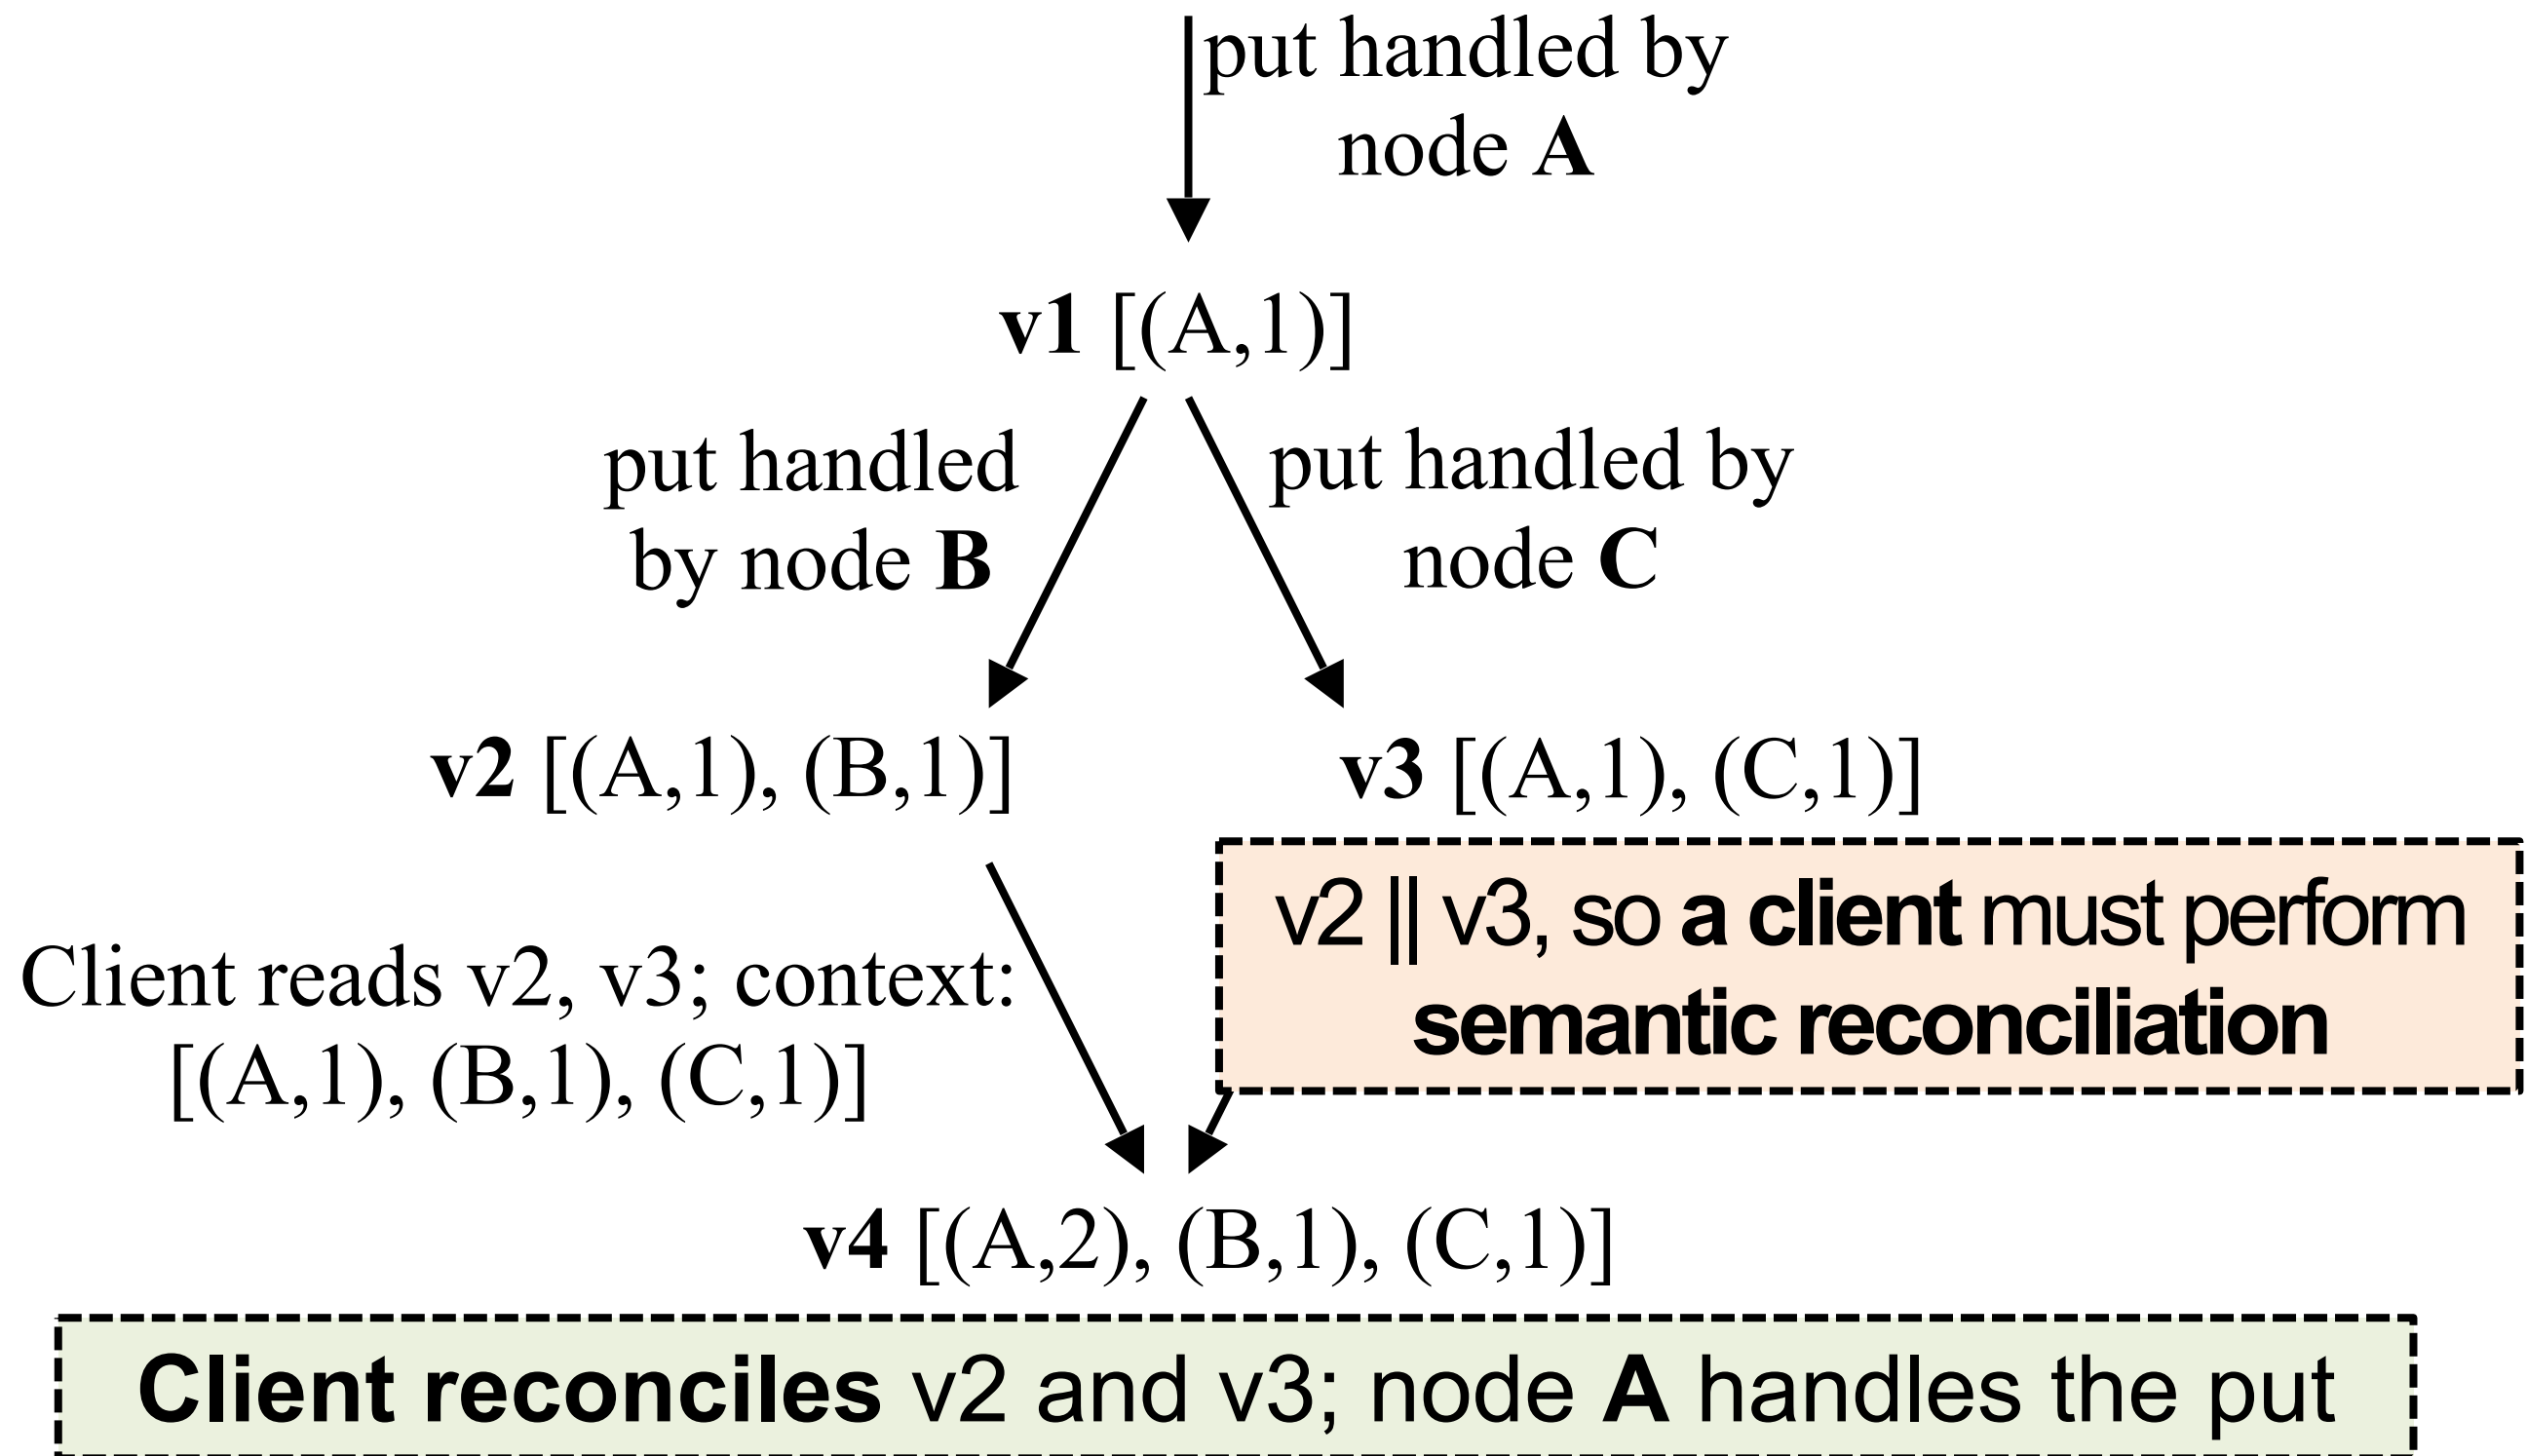
- Then users **must supply that context** to put()s that modify the same key

Version vectors (auto-resolving case)



$v2 > v1$, so Dynamo nodes **automatically drop $v1$** , for $v2$

Version vectors (app-resolving case)



Removing threats to durability

Hinted handoff node **crashes before it can replicate data** to node in **preference list**

- Need another way to **ensure** that each key-value pair is **replicated N times**

Mechanism: replica synchronization

- Nodes nearby on ring periodically **gossip**
 - » **Compare** the (k, v) pairs they hold
 - » **Copy** any missing keys the other has

How to **compare and copy** replica state **quickly and efficiently?**

Efficient synchronization with Merkle trees

Merkle trees hierarchically summarize the key-value pairs a node holds

One Merkle tree for each **virtual node key range**

- **Leaf node** = hash of **one key's value**
- **Internal node** = hash of **concatenation of children**

Compare roots; if match, values match


- If they **don't match**, compare **children**
 - » **Iterate** this process down the tree

Merkle tree reconciliation


B is missing orange key; **A** is missing green one

Exchange and compare hash nodes from root downwards, **pruning when hashes match**

A's values:



B's values:



Finds differing keys **quickly** and with minimum information exchange

Dynamo: Take-away ideas

Consistent hashing broadly useful for replication—not only in P2P systems

Extreme emphasis on **availability and low latency**, unusually, at the **cost of some inconsistency**

Eventual consistency lets writes and reads return quickly, **even when partitions and failures**

Version vectors allow some **conflicts to be resolved** automatically; others left to application