

# Issues

---

- ▶ Also, subject to quirks of the architecture
  - ▶ Example: x86
  - ▶ fails silently if some privileged instructions execute without privilege
  - ▶ doesn't have clean separation between privileged and non privileged instructions
  - ▶ Backward compatibility means difficult to improve

# Issues

---

- ▶ Consider Intel x86 popf instruction
  - ▶ Loads CPU flags register from contents of the stack
  - ▶ If CPU in privileged mode -> all flags replaced
  - ▶ If CPU in user mode -> on some flags replaced
  - ▶ No trap is generated

# Binary translation

---

- ▶ Use binary translation to modify OS to rewrite silent failure instructions
- ▶ More aggressive translation can be used
  - ▶ Translate OS mode instructions to equivalent VMM instructions
  - ▶ Some operations still expensive
    - Cache for future use
    - Used by VMWare ESXi and Microsoft Virtual Server
- ▶ Performance on x86 typically ~80-95% of native

# Paravirtualization

---

- ▶ Modify the OS to make it aware of the hypervisor
  - ▶ Can avoid the tricky features  
Aware of the fact it is virtualized
  - ▶ Can implement optimizations
- ▶ Comparison to binary translation?
- ▶ Amount of code change?
  - 1.36% of Linux, 0.04% for Windows

# Hardware assistance

---

- ▶ All virtualization needs some HW support
  - ▶ More support -> more feature rich, stable, better performance of guests
- ▶ Intel added new VT-x instructions in 2005 and AMD the AMD-V instructions in 2006
  - ▶ CPUs with these instructions remove need for binary translation
  - ▶ Generally define more CPU modes – “guest” and “host”
  - ▶ VMM can enable host mode, define characteristics of each guest VM, switch to guest mode and guest(s) on CPU(s)
  - ▶ In guest mode, guest OS thinks it is running natively, sees devices (as defined by VMM for that guest)
  - ▶ Access to virtualized device, privileged instructions cause trap to VMM
  - ▶ CPU maintains VCPU, context switches it as needed
  - ▶ HW support for Nested Page Tables, DMA, interrupts as well over time

Next: Virtualization 2