

CSCI 381/780

Cloud Computing

Virtualization 2: Container

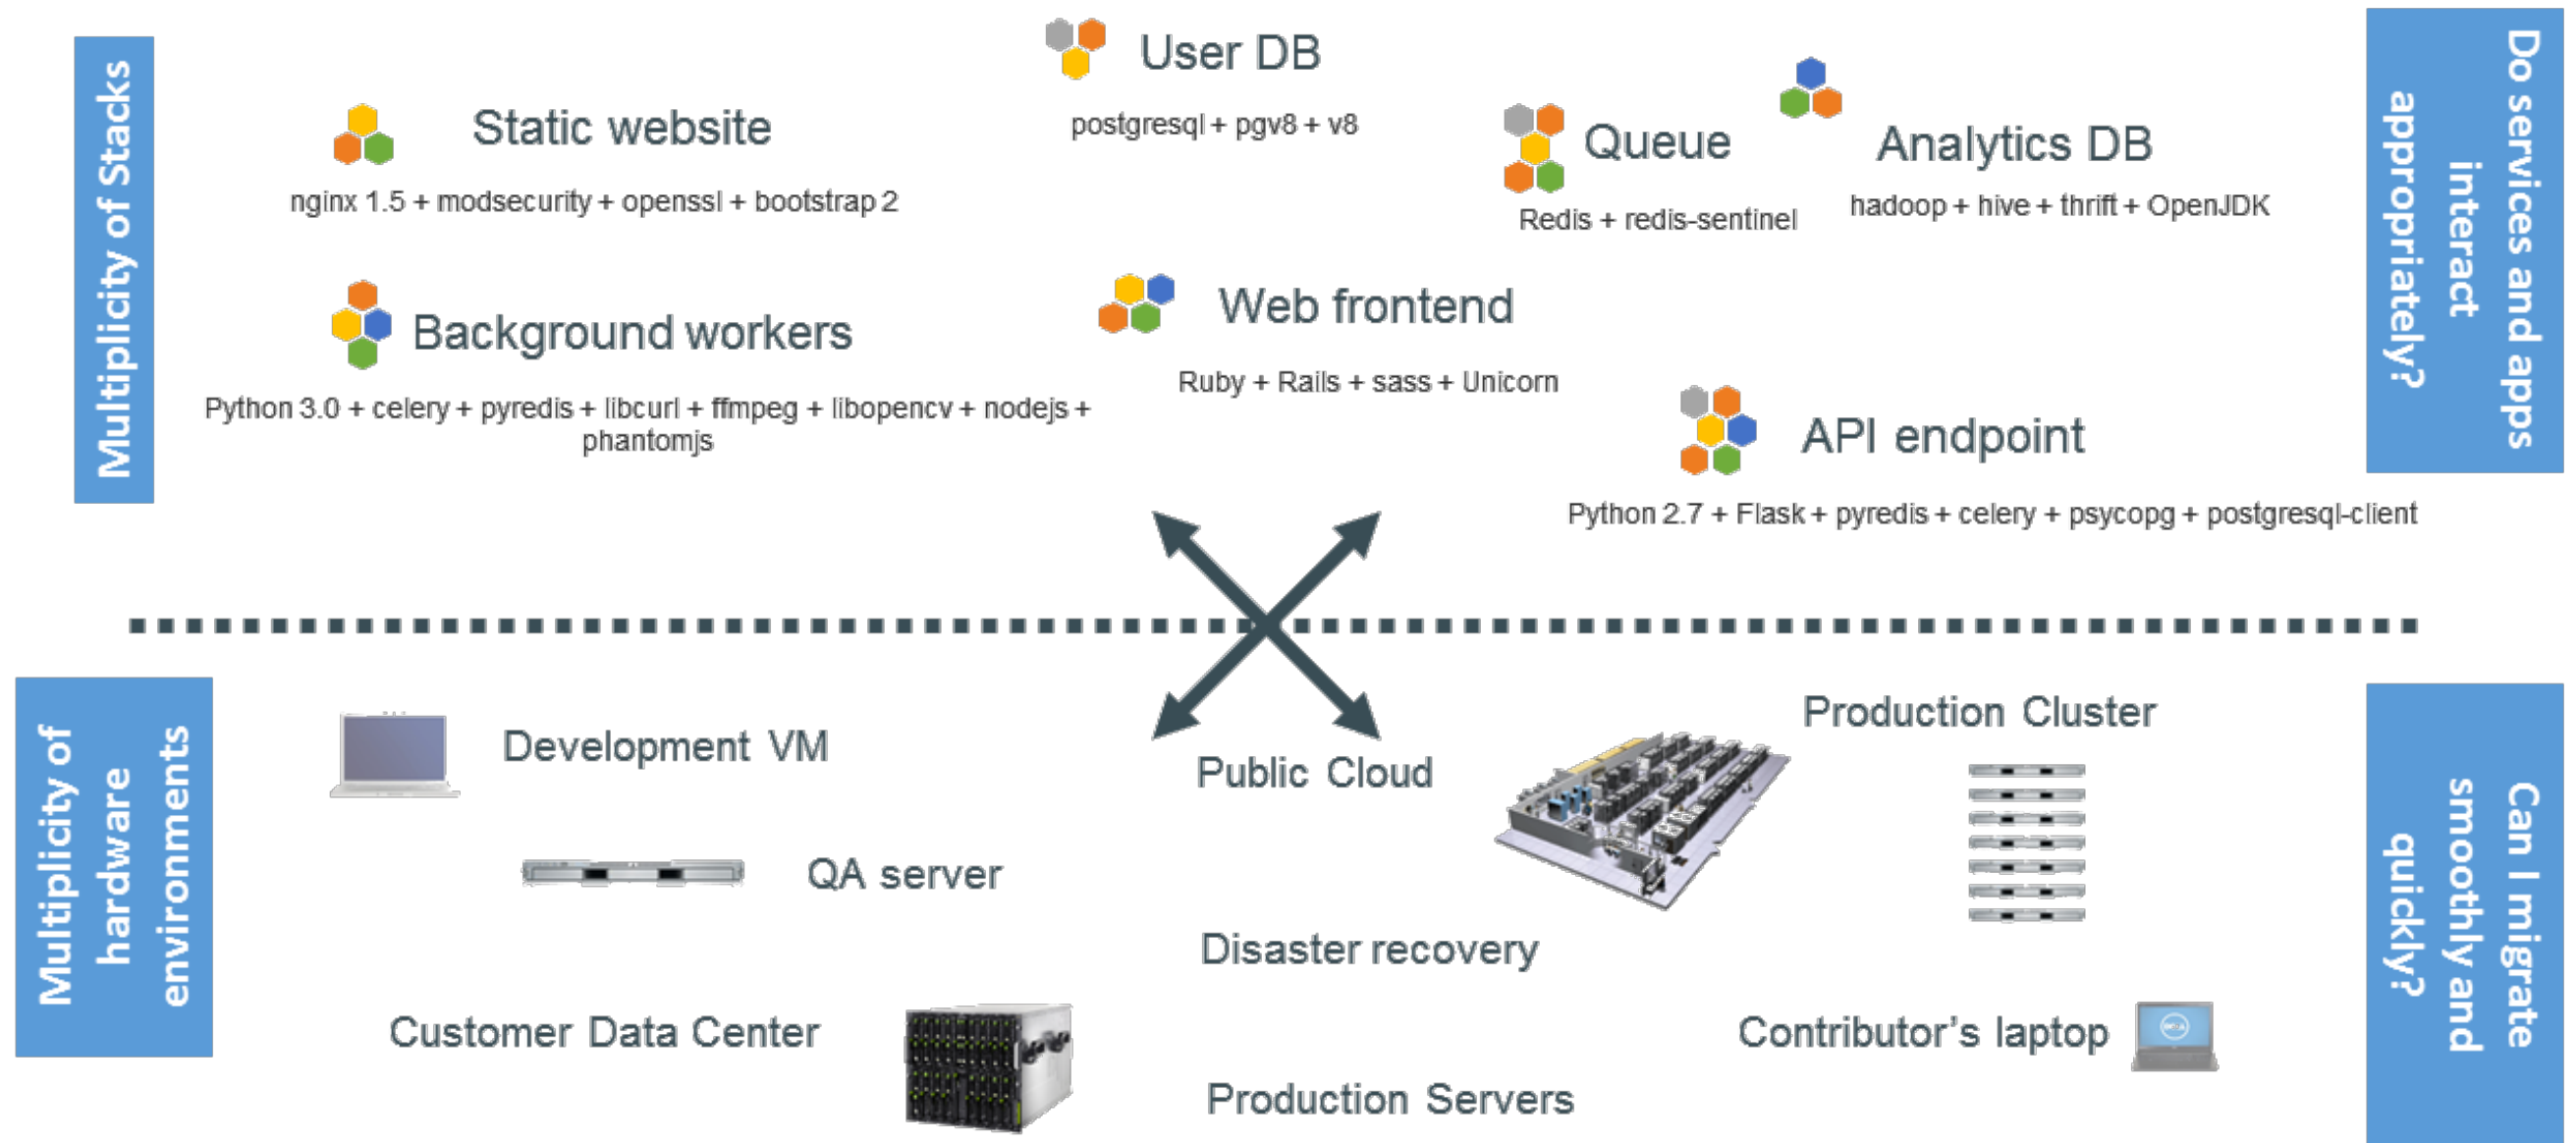
Jun Li
Queens College
















Are VMs Fit for (All) Today's (Cloud) Usages?

- ▶ Performance overhead of indirections (guest OS and hypervisor)
- ▶ Large memory footprint
- ▶ Slow startup time
- ▶ License and maintenance cost of guest OS
- ▶ Do we really need to virtualize hardware and a full OS?
- ▶ What about DevOps?


The Challenge



The Matrix from Hell

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

Cargo Transportation Pre-1960



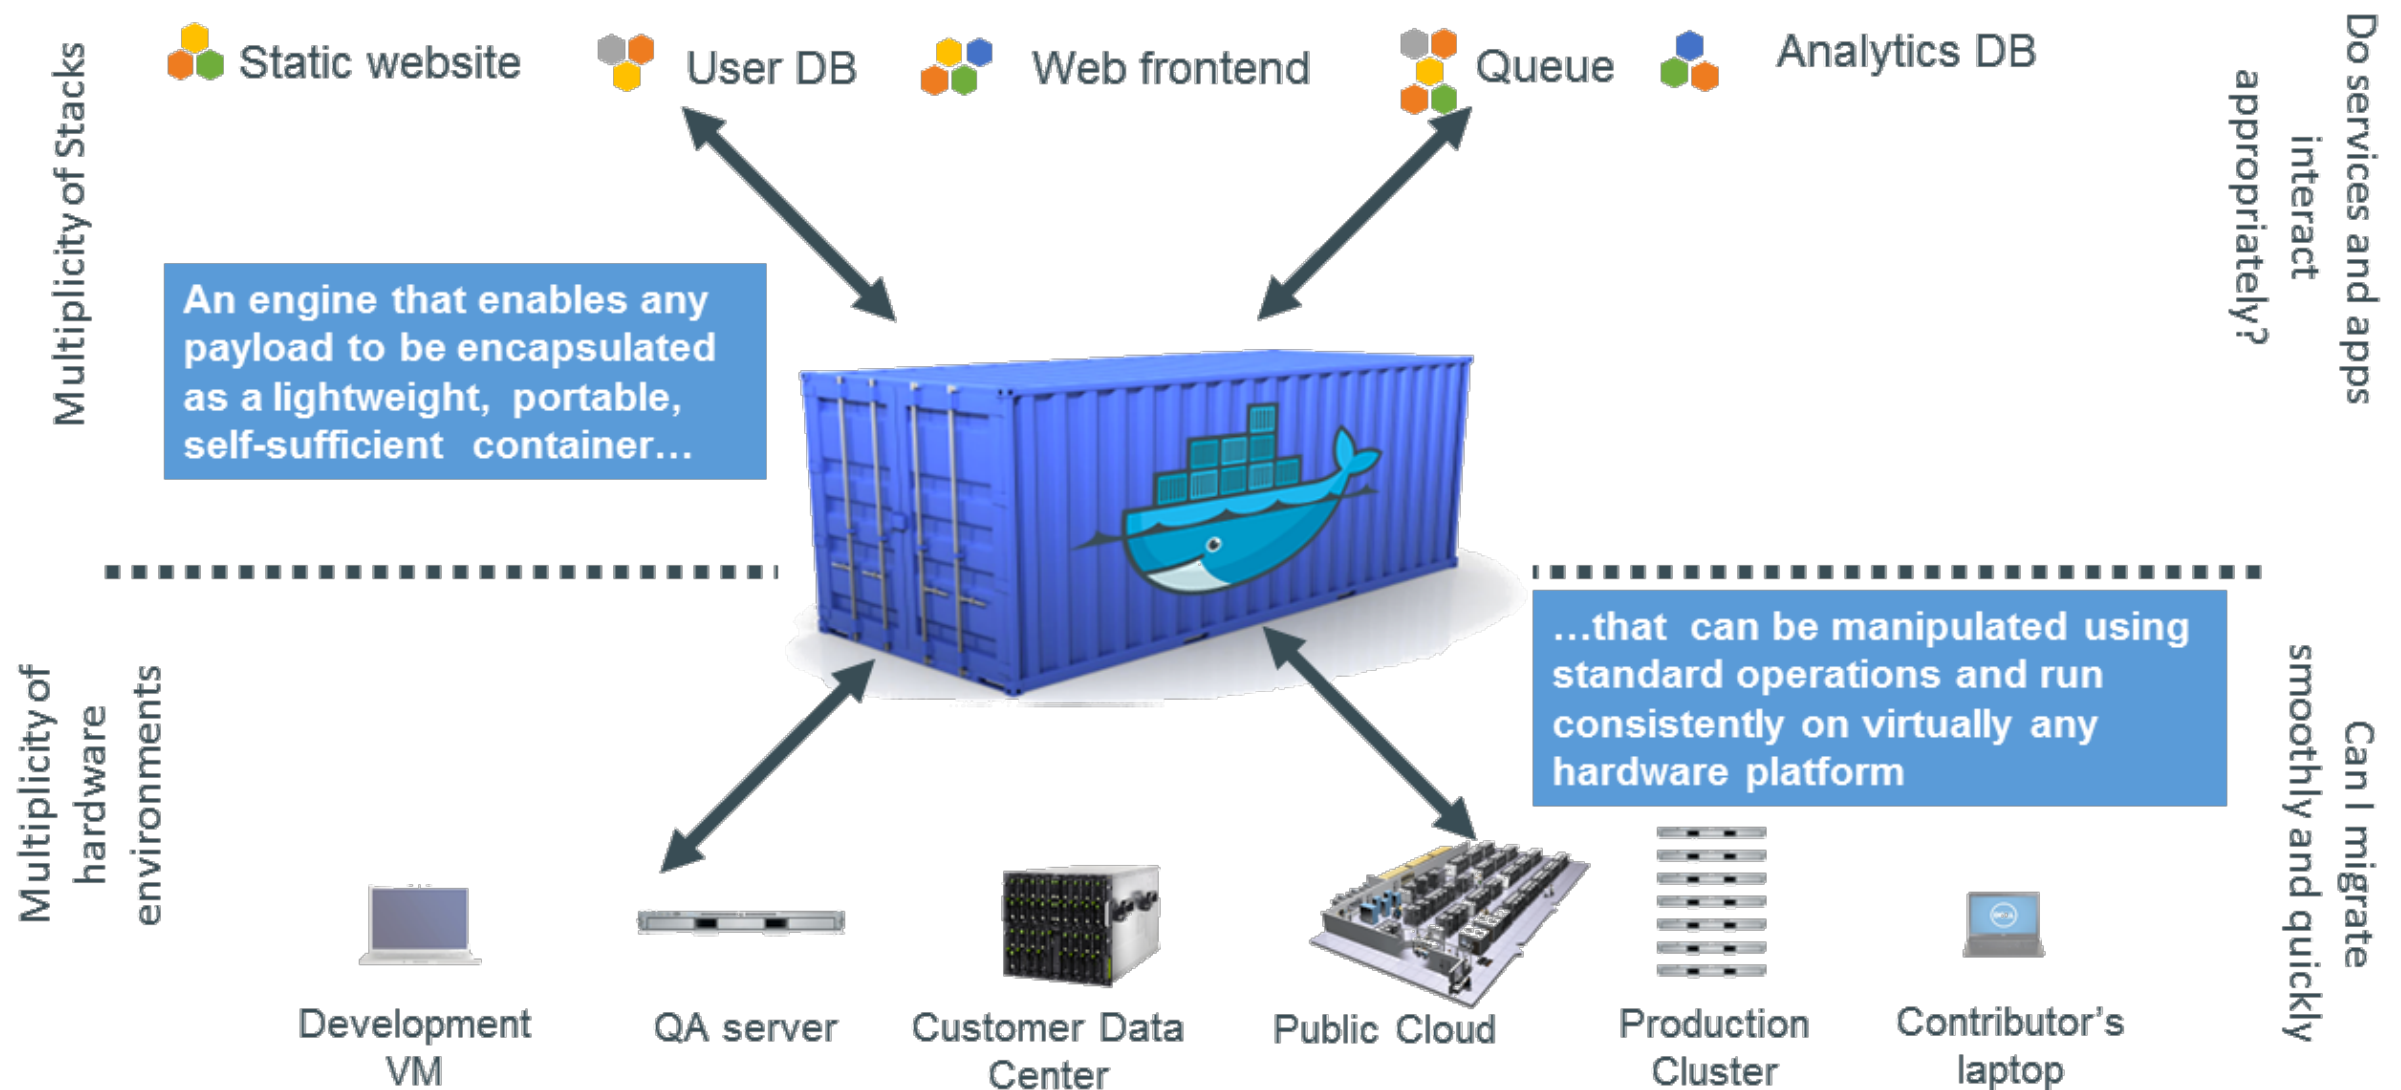
Also a Matrix from Hell

	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
							

Solution: Shipping Container



Docker: Container for Code




	Static website							
	Web frontend							
	Background workers							
	User DB							
	Analytics DB							
	Queue							
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers



Why Does It Work? Separation of Concerns

- **Dan the Developer**

- Worries about what's "inside" the container
 - His code
 - His Libraries
 - His Package Manager
 - His Apps
 - His Data
- All Linux servers look the same

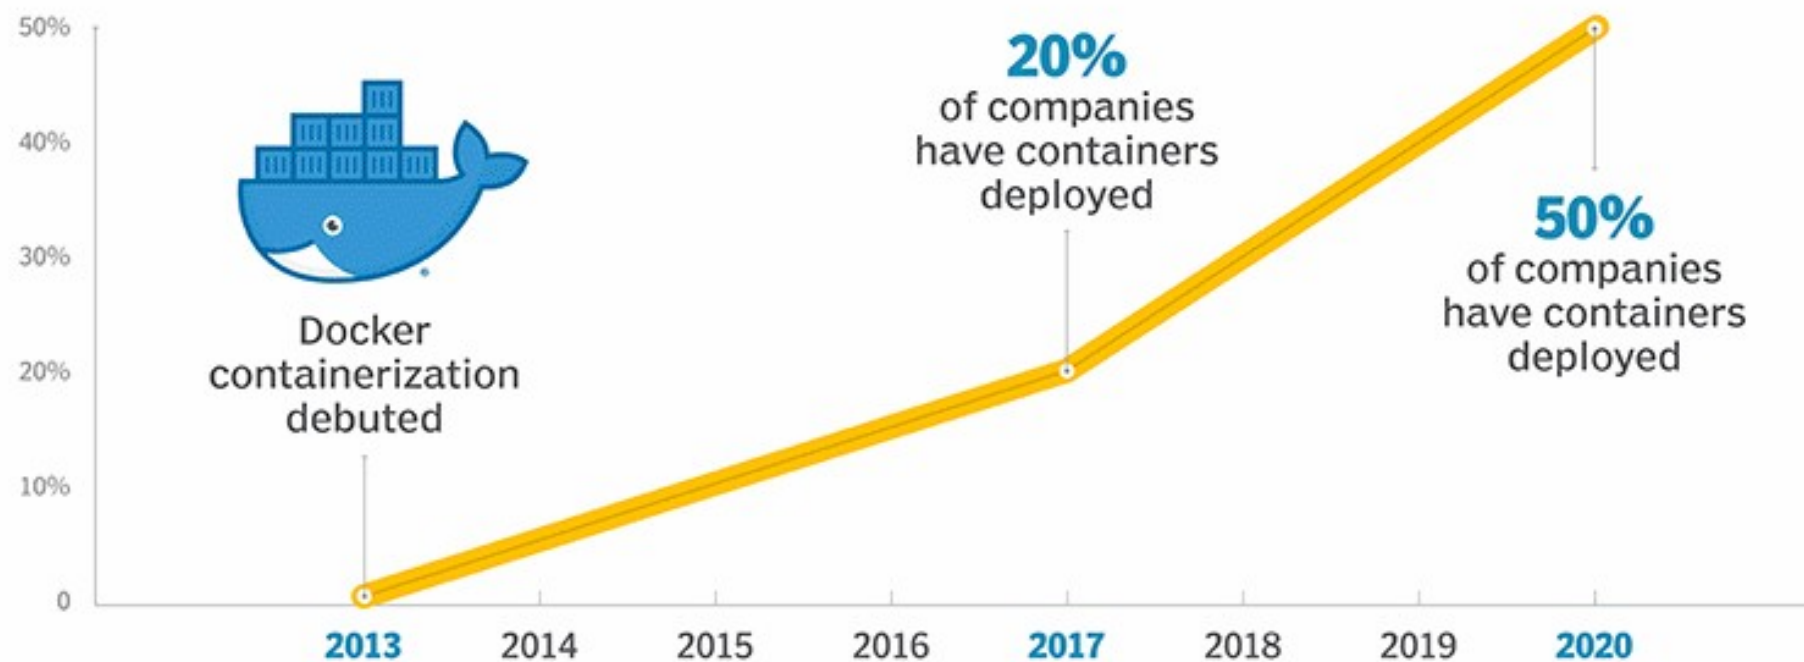


Major components of the container:

- **Oscar the Ops Guy**

- Worries about what's "outside" the container
 - Logging
 - Remote access
 - Monitoring
 - Network config
- All containers start, stop, copy, attach, migrate, etc. the same way

Containerization timeline



SOURCE: GARTNER

Linux Containers

- ▶ Run everywhere
 - ▶ Regardless of kernel version
 - ▶ Regardless of host distro
 - ▶ Physical or virtual, cloud or not
 - ▶ Container and host architecture must match...
- ▶ Run anything
 - ▶ If it can run on the host, it can run in the container
 - ▶ If it can run on a Linux kernel, it can run, it can run in the container

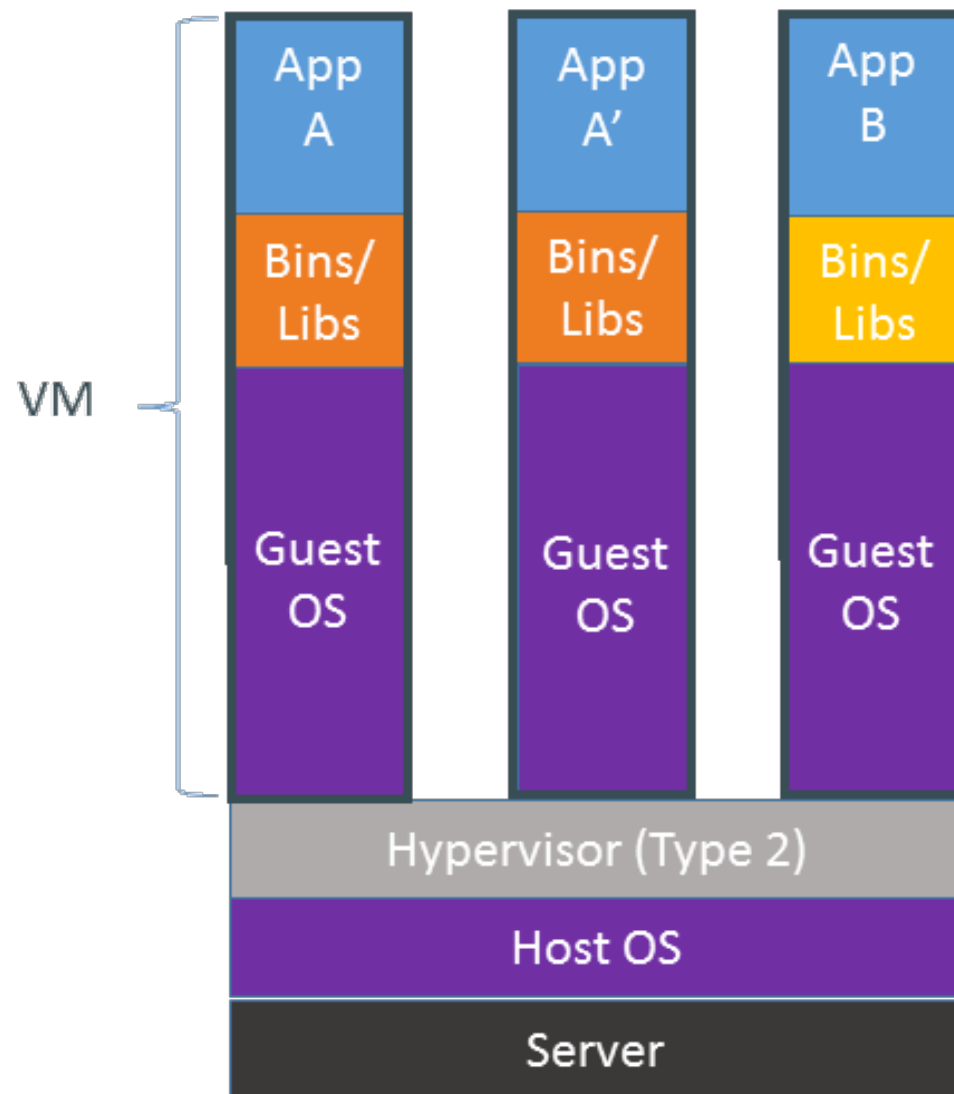
At High-Level: It Looks Like a VM

- ▶ Own process space
- ▶ Own network interface
- ▶ Can run stuff as root
- ▶ Can have its own /sbin/init (different from the host)

At Low-Level: OS-Level Virtualization


- ▶ Containers run on a host OS directly (and share the OS)
- ▶ Run as processes
- ▶ OS provides resource isolation and namespace isolation

VM vs Container



Containers are isolated, but share OS and, where appropriate, bins/libraries

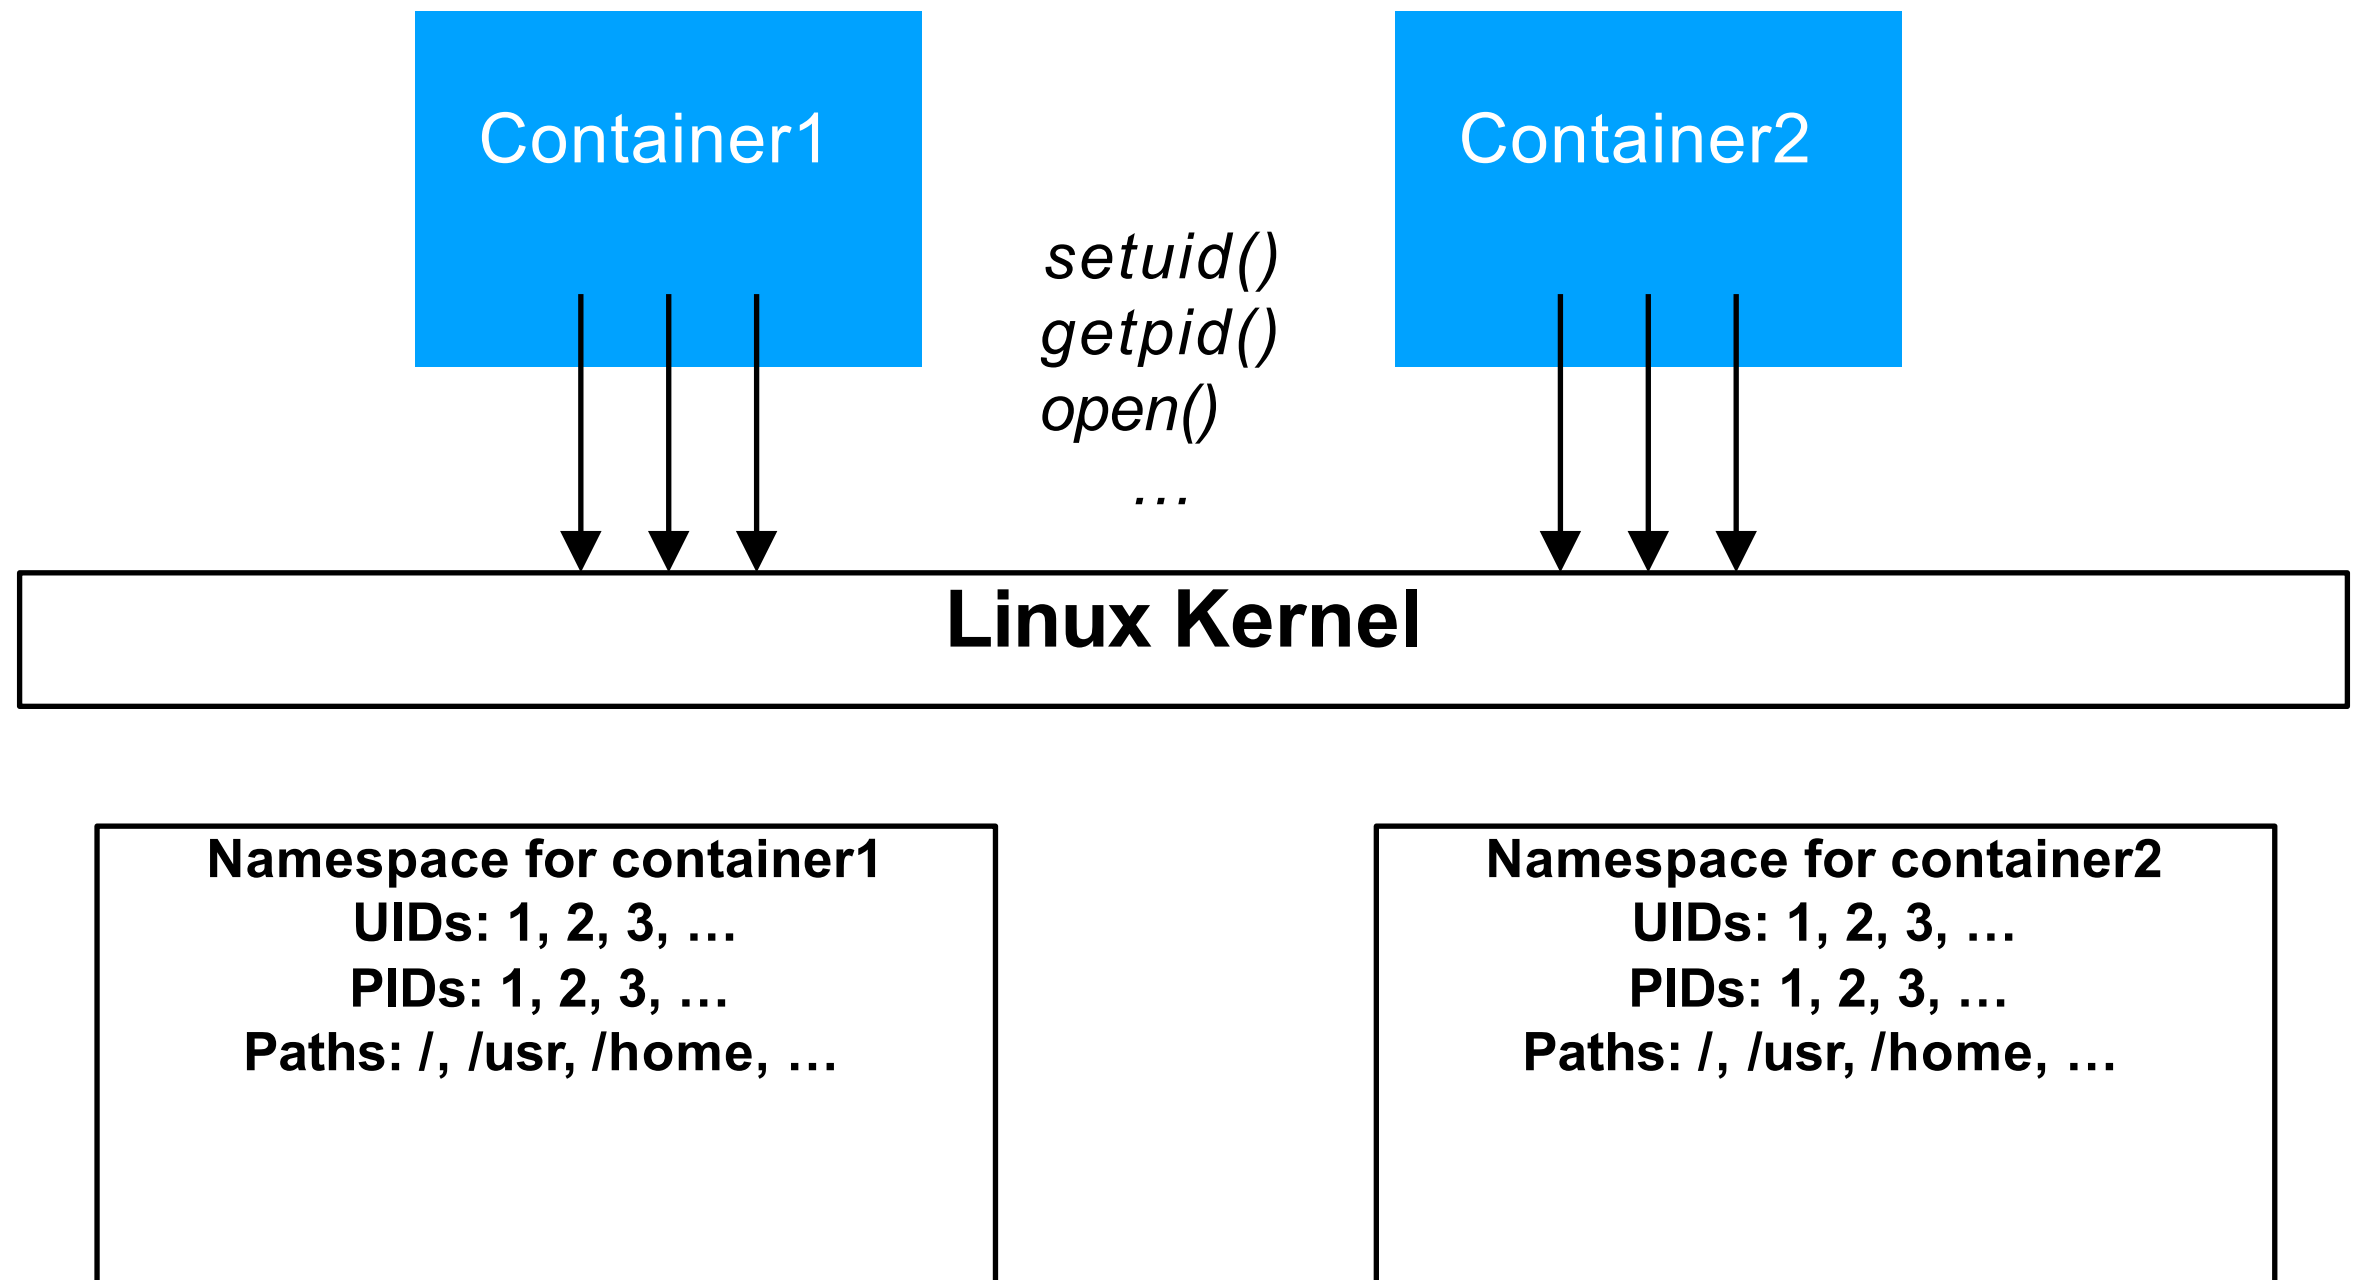
...result is significantly faster deployment, much less overhead, easier migration, faster restart



Using Namespaces to Separate “Views” of Users

- ▶ Namespace: naming domain for various resources
 - ▶ User IDs (UIDs)
 - ▶ Process IDs (PIDs)
 - ▶ File paths (mnt)
 - ▶ Network sockets
 - ▶ Pipe names

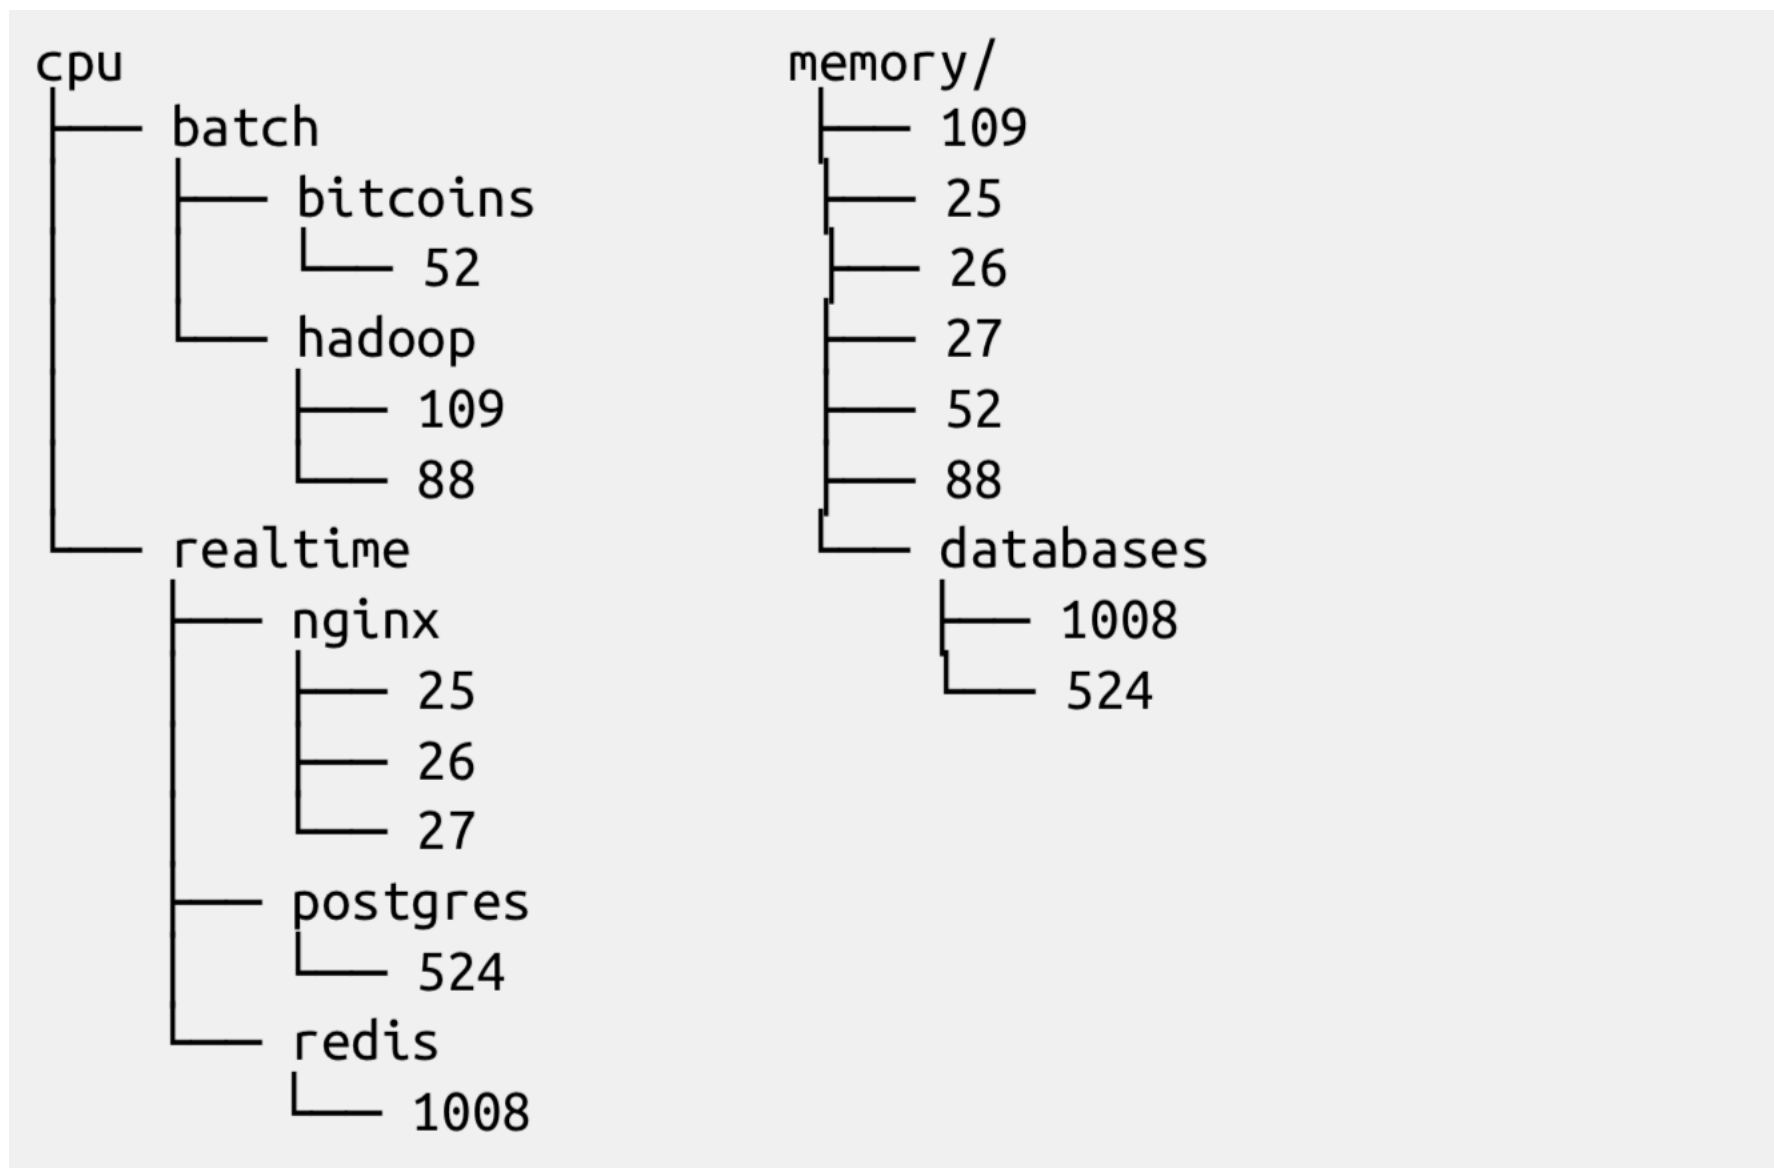
Namespaces Isolated by Kernel



Isolating Resources with cgroups

- ▶ Linux Control Groups (cgroups): collection of Linux processes
 - ▶ Limits resource usages at group level (e.g., memory, CPU, device)
 - ▶ Fair sharing of resources
 - ▶ Track resource utilization (e.g., could be used for billing/management)
 - ▶ Control processes (e.g., pause/resume, checkpoint/restore)

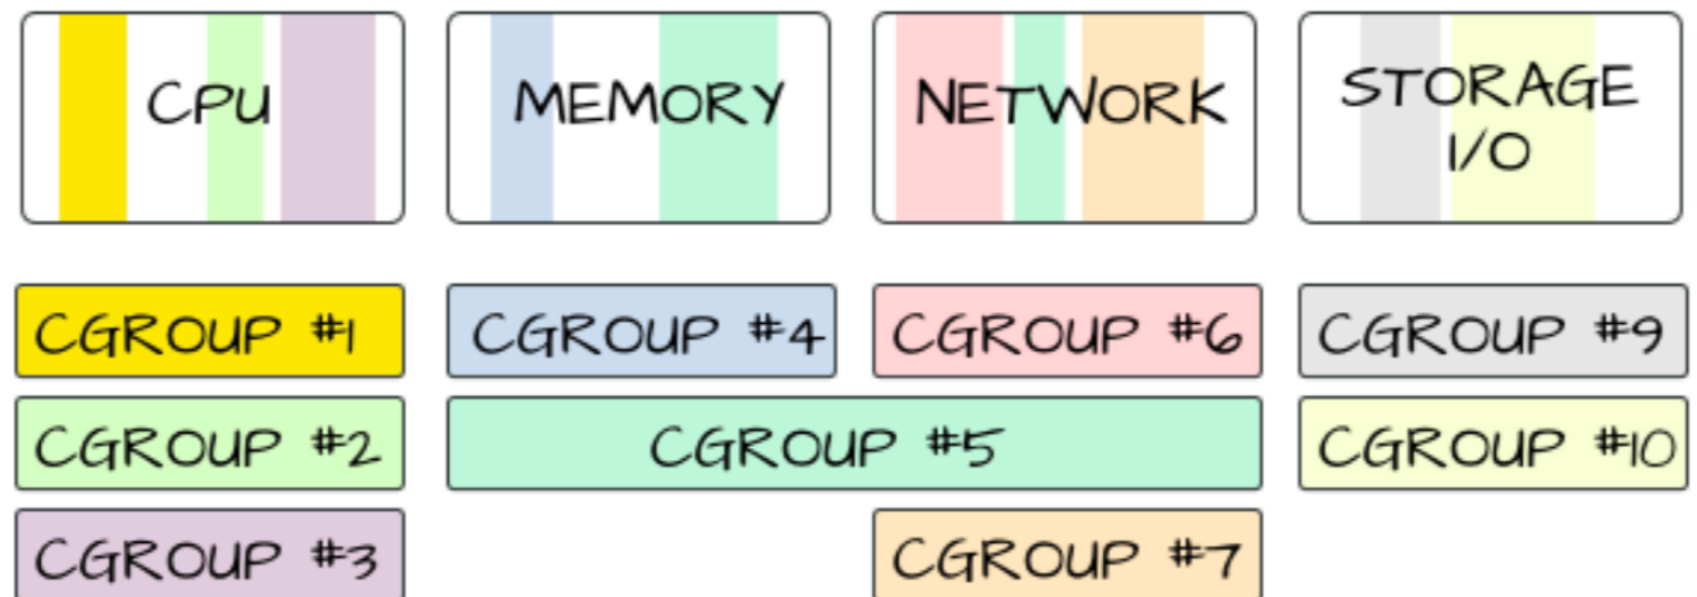
Example



Resource Isolation

Cgroups : Isolation and accounting

- cpu
- memory
- block i/o
- devices
- network
- numa
- freezer

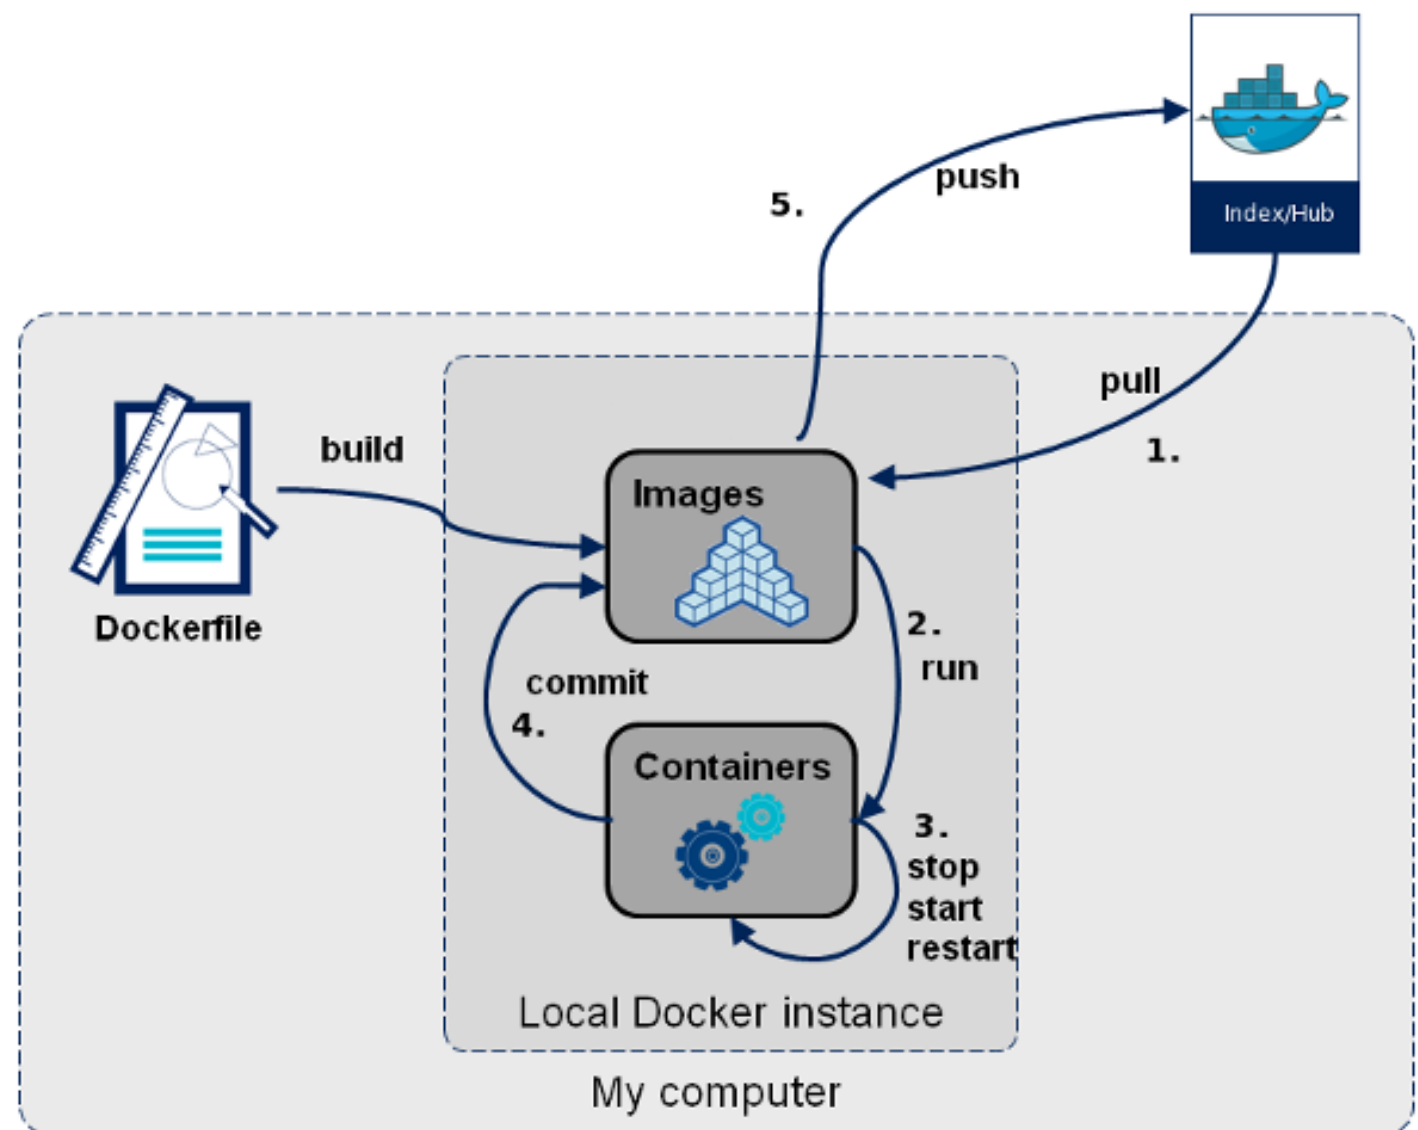


Efficiency: *almost* no overhead


- ▶ Processes are isolated, but run straight on the host
- ▶ CPU performance = native performance
- ▶ Memory performance = a few % shaved off for (optional) accounting
- ▶ Network performance = small overhead; can be optimized to zero overhead

Docker

- ▶ Docker is a platform for developing, shipping & running application using container based virtualization technology.
- ▶ founded as dotCloud in 2010
- ▶ Docker: A container engine written in Go
- ▶ popularized containerization
- ▶ Provides a straightforward way to build and run containers




Docker Adoption Behavior



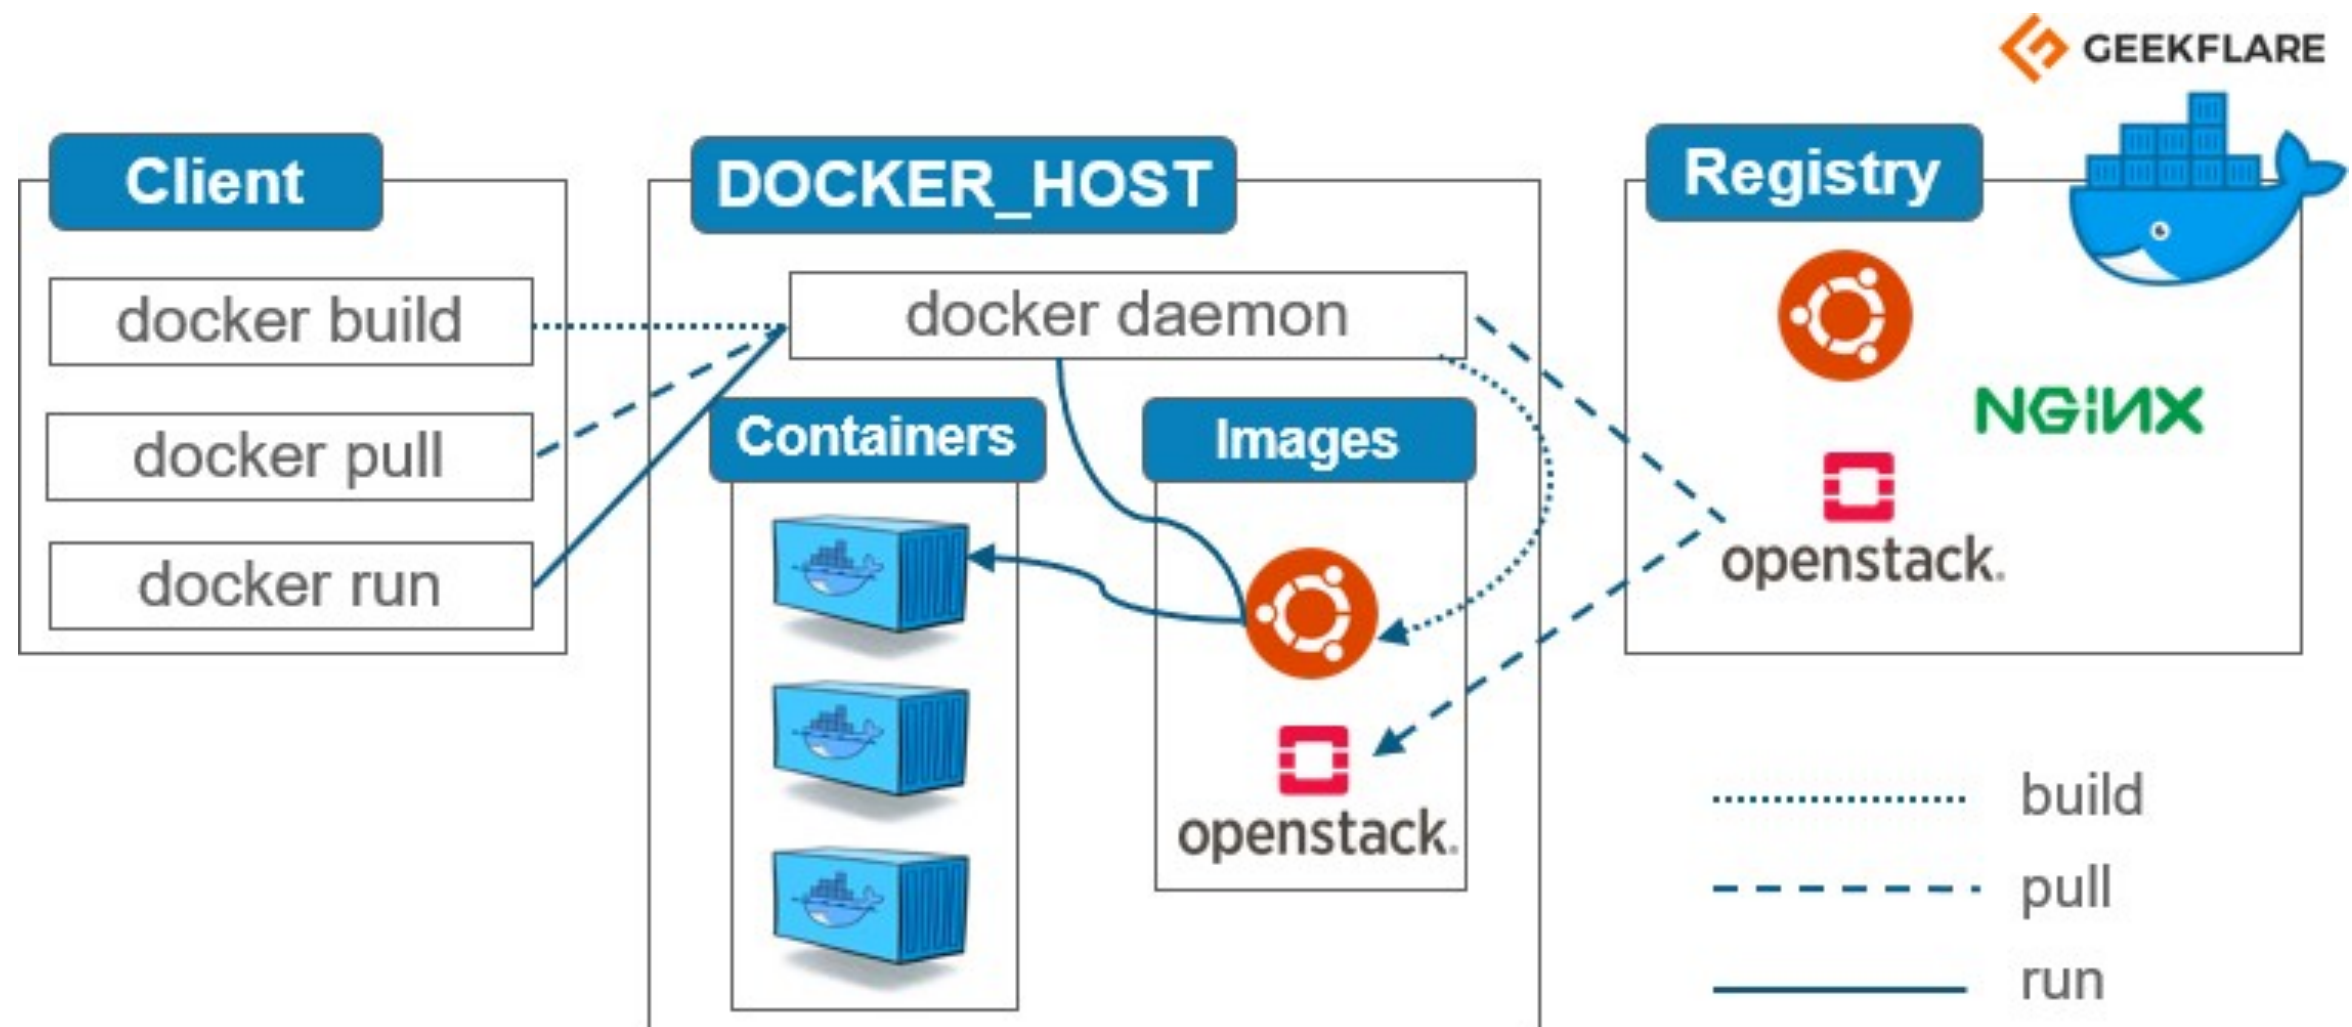
Source: Datadog

Portion of Hosts Running Docker



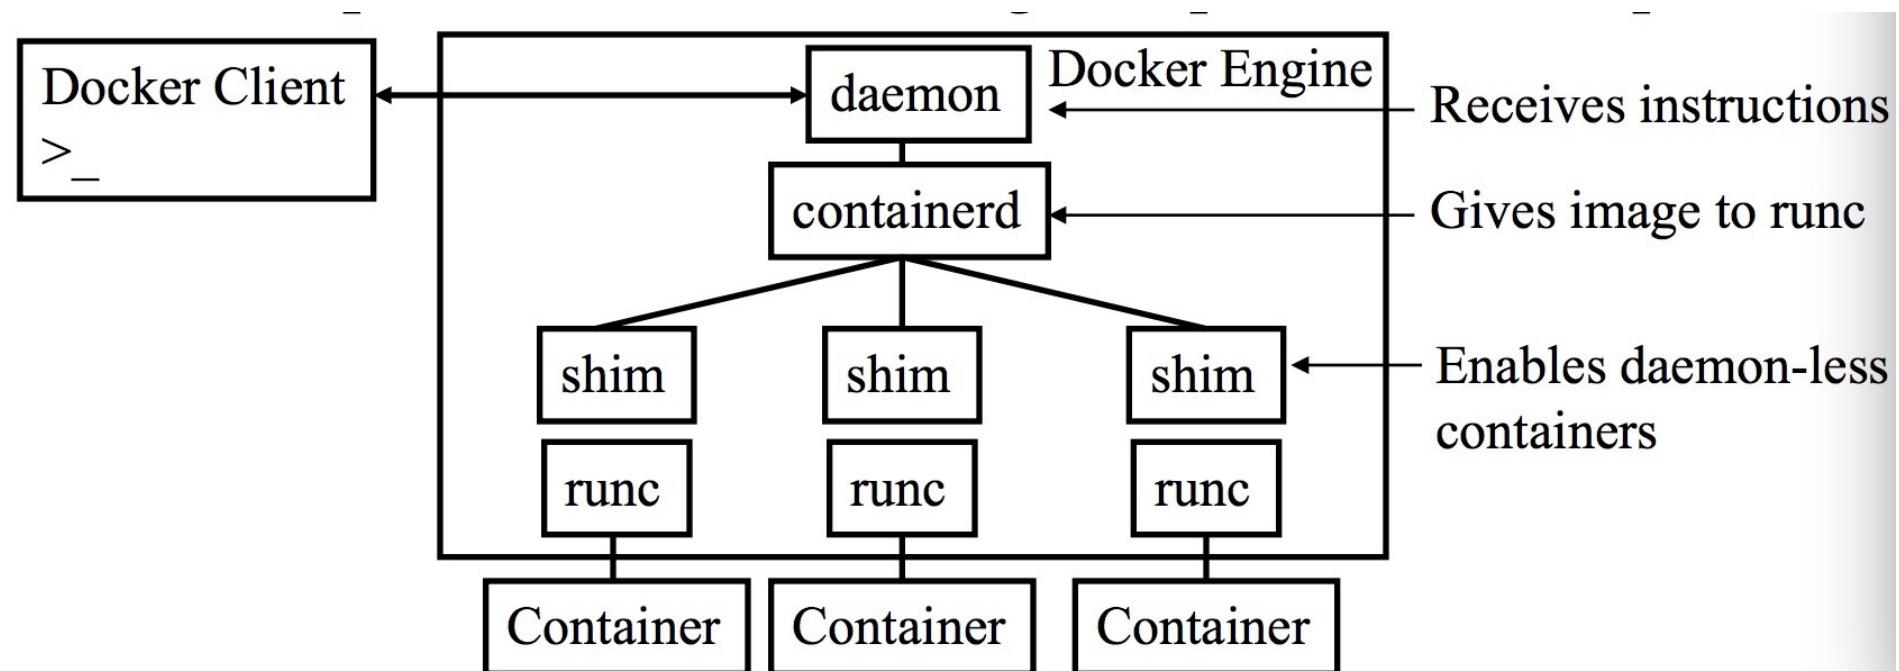
Source: Datadog

Docker Architecture




Docker Engine

- daemon: Rest API (receiving instructions) and other features
- containerd: Execution logic (e.g., start, stop, pause, unpause, delete containers)
- runc: A lightweight runtime CLI



Docker Image


- ▶ A Docker image is a binary that includes all of the requirements for running a single Docker container, as well as metadata describing its needs and capabilities.
- ▶ Each image consists of a series of layers using the union file system
- ▶ When you change an image a new layer is created.



JULIA EVANS
@b0rk

overlay filesystems

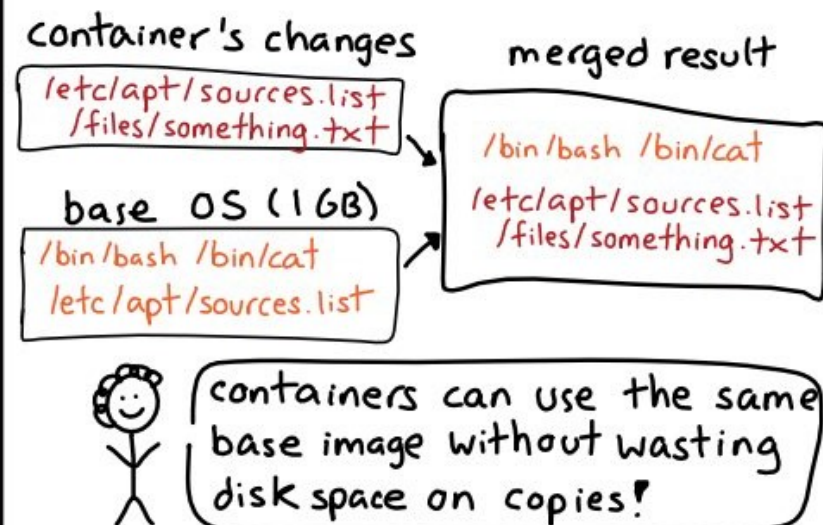
Docker images can be really big



every container needs a "copy" of its image



solution: ★ overlays ★



how to overlay

Linux has an 'overlayfs' driver that you can use to overlay directories like this:

```
$ mount -t overlay overlay  
-o lowerdir=/lower,upperdir=/upper,workdir=/work  
/merged
```

↑ base directory, will be read only

↑ where changes will go

↑ must be empty

Try it out! It's really easy.

how Docker runs containers

- ① unpack the base image into a directory
- ② make an empty directory for changes
- ③ overlay the new directory on top of the base
- ④ start the container!

Building an Image

- ▶ A Dockerfile that contains a set of instructions that tell Docker how to build our image. The Dockerfile can be used to generate an image stored on your local system.
- ▶ Docker daemon does actual build process. Contents of current context (folder) is sent to the daemon.
- ▶ Each Instruction creates a new layer in the image. The cache from previous builds is used as much as possible.

Dockerfile

```
FROM tomcat:7.0.62-jre8
MAINTAINER Jeff Ellin jeff.ellin@jeffellin.com
ENV CORE_SQL_URL "jdbc:postgresql://localhost:5432/core"
ENV CORE_SQL_USERNAME "tamr"
ENV CORE_SQL_PASSWORD "12345"
#Enable use of gui admin tool
add tomcat-users.xml $CATALINA_HOME/conf
#add the tamr war

add tamr.war /tamr/tamr.war
add catalina.sh $CATALINA_HOME/bin
RUN mv /tamr/*.war $CATALINA_HOME/webapps
```

base image

environment variables

local files

execute command