

# CSCI 381/780

## Cloud Computing

### Streaming Analytics

---

Jun Li  
Queens College



# MapReduce

---

- ▶ Batch Processing => Need to wait for entire computation on large dataset to complete
- ▶ Not intended for long-running stream-processing

# Challenges

---

- ▶ Large amounts of data => Need for real-time views of data
  - ▶ Social network trends, e.g., Twitter real-time search
  - ▶ Website statistics, e.g., Google Analytics
  - ▶ Intrusion detection systems, e.g., in most data centers
- ▶ Process large amounts of data
  - ▶ With latencies of few seconds
  - ▶ With high throughput

# Storm

---

- ▶ Apache Project
- ▶ <http://storm.apache.org/>
- ▶ Highly active JVM project
- ▶ Multiple languages supported via API
  - ▶ Python, Ruby, etc.
- ▶ Used by over 30 companies including
  - ▶ Twitter: For personalization, search
  - ▶ Flipboard: For generating custom feeds
  - ▶ Weather Channel, WebMD, etc.

# Tuple

---

- ▶ An ordered list of elements
- ▶ E.g., < tweeter, tweet >
  - ▶ E.g., < “Miley Cyrus”, “Hey! Here’s my new song!” >
  - ▶ E.g., < “Justin Bieber”, “Hey! Here’s MY new song!” >
- ▶ E.g., < URL, clicker-IP, date, time >
  - ▶ E.g., < coursera.org, 101.102.103.104, 4/4/2014, 10:35:40 >
  - ▶ E.g., < coursera.org, 101.102.103.105, 4/4/2014, 10:35:42 >


Tuple



# Stream

---


- ▶ Sequence of tuples
  - ▶ Potentially unbounded in number of tuples
- ▶ Social network example:
  - ▶ <“Miley Cyrus”, “Hey! Here’s my new song!”>,  
    <“Justin Bieber”, “Hey! Here’s MY new song!”>,  
    <“Rolling Stones”, “Hey! Here’s my old song that’s still a super-hit!”>, ...
- ▶ Website example:
  - ▶ <coursera.org, 101.102.103.104, 4/4/2014, 10:35:40>, <coursera.org, 101.102.103.105, 4/4/2014, 10:35:42>, ...



# Spout

---


- ▶ A Storm entity (process) that is a source of streams
- ▶ Often reads from a crawler or DB



# Bolt

---

- ▶ A Storm entity (process) that
  - ▶ Processes input streams
  - ▶ Outputs more streams for other bolts






# Topology

---


- ▶ A directed graph of spouts and bolts (and output bolts)
- ▶ Corresponds to a Storm “application”



# Topology

---

- Can have cycles if the application requires it



# Bolts come in many Flavors

---

- ▶ Operations that can be performed
  - ▶ Filter: forward only tuples which satisfy a condition
  - ▶ Joins: When receiving two streams A and B, output all pairs (A,B) which satisfy a condition
  - ▶ Apply/transform: Modify each tuple according to a function
  - ▶ And many others
- ▶ But bolts need to process a lot of data
  - ▶ Need to make them fast



# Parallelizing Bolts

---

- ▶ Have multiple processes (“tasks”) constitute a bolt
- ▶ Incoming streams split among the tasks
- ▶ Typically each incoming tuple goes to one task in the bolt
  - ▶ Decided by “Grouping strategy”
- ▶ Three types of grouping are popular

# Storm Task

- ▶ Spouts and bolts execute as many tasks across the cluster
- ▶ When a tuple is emitted, which task does it go to? User programmable:
  - ▶ Shuffle grouping: pick a random task
  - ▶ Fields grouping: consistent hashing on a subset of tuple fields
  - ▶ All grouping: send to all tasks
  - ▶ Global grouping: pick task with lowest id



# Storm Cluster

---

- ▶ Master node
  - ▶ Runs a daemon called *Nimbus*
  - ▶ Responsible for
    - ▶ Distributing code around cluster
    - ▶ Assigning tasks to machines
    - ▶ Monitoring for failures of machines
- ▶ Worker node
  - ▶ Runs on a machine (server)
  - ▶ Runs a daemon called *Supervisor*
  - ▶ Listens for work assigned to its machines
  - ▶ Runs “Executors”(which contain groups of tasks)
- ▶ Zookeeper
  - ▶ Coordinates Nimbus and Supervisors communication
  - ▶ All state of Supervisor and Nimbus is kept here

# Failures

---

- ▶ A tuple is considered failed when its topology (graph) of resulting tuples fails to be fully processed within a specified timeout
- ▶ Anchoring: Anchor an output to one or more input tuples
  - ▶ Failure of one tuple causes one or more tuples to replayed

# API For Fault-Tolerance (OutputCollector)


---

- ▶ `Emit(tuple, output)`
  - ▶ Emits an output tuple, perhaps anchored on an input tuple (first argument)
- ▶ `Ack(tuple)`
  - ▶ Acknowledge that you (bolt) finished processing a tuple
- ▶ `Fail(tuple)`
  - ▶ When a tuple fails, Storm will attempt to replay it from its source and reprocess it.
  - ▶ If the tuple cannot be replayed or reprocessed, Storm will consider it a permanent failure and remove it from its tracking data structures.
- ▶ Must remember to ack/fail each tuple
  - ▶ Each tuple consumes memory. Failure to do so results in memory leaks.



# Stateful Stream Processing

- ▶ Traditional streaming systems have a **record-at-a-time** processing model
  - ▶ Each node has mutable state
  - ▶ For each record, update state and send new records



- State is lost if node dies!
- Making stateful stream processing be fault-tolerant is challenging

# Existing Streaming Systems

---

- ▶ Storm
  - ▶ Replays record if not processed by a node
  - ▶ Processes each record *at least once*
  - ▶ May update mutable state twice!
  - ▶ Mutable state can be lost due to failure!

# What is Spark Streaming?

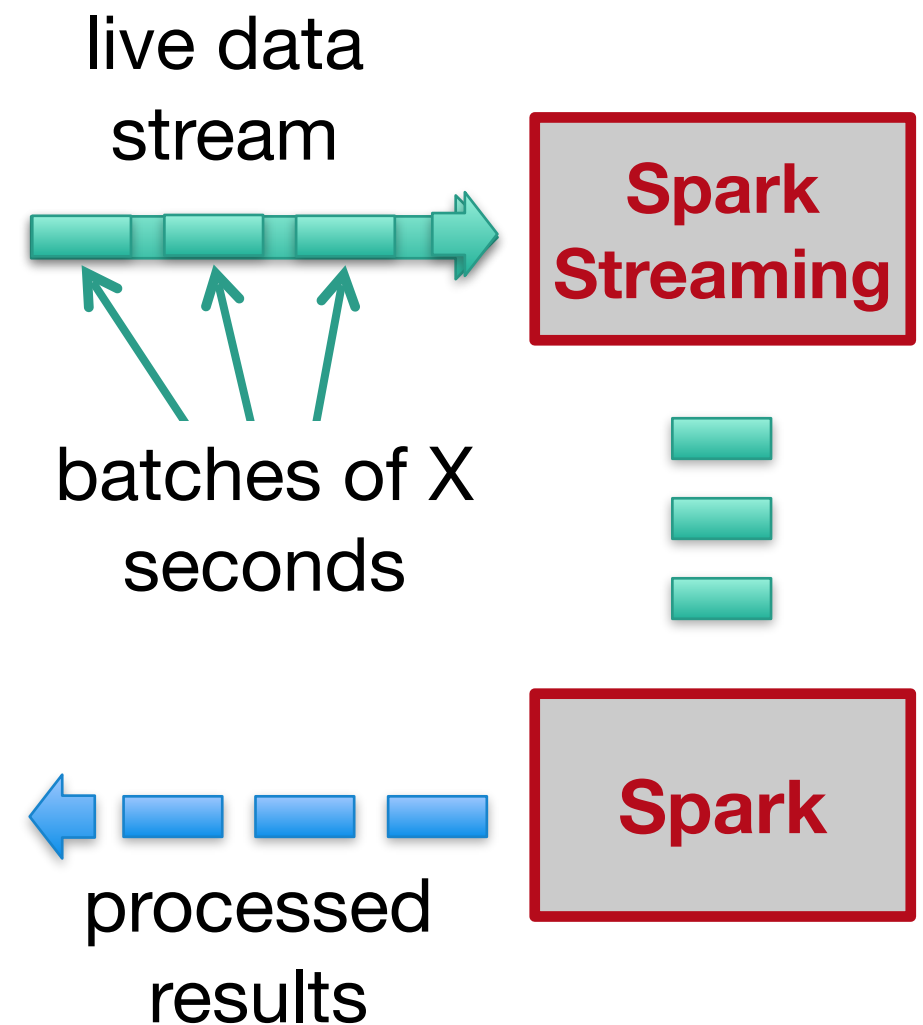
---

- ▶ Extends Spark for doing large scale stream processing
- ▶ Scales to 100s of nodes and achieves second scale latencies
- ▶ Efficient and fault-tolerant stateful stream processing
- ▶ Simple batch-like API for implementing complex algorithms

# Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**


- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches



# Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**

- Batch sizes as low as 1/2 second, latency of about 1 second
- Potential for combining batch processing and streaming processing in the same system



# Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()
```


**DStream:** a sequence of RDDs representing a stream of data

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream



stored in memory as an  
RDD (immutable,  
distributed)

# Example – Get hashtags from Twitter

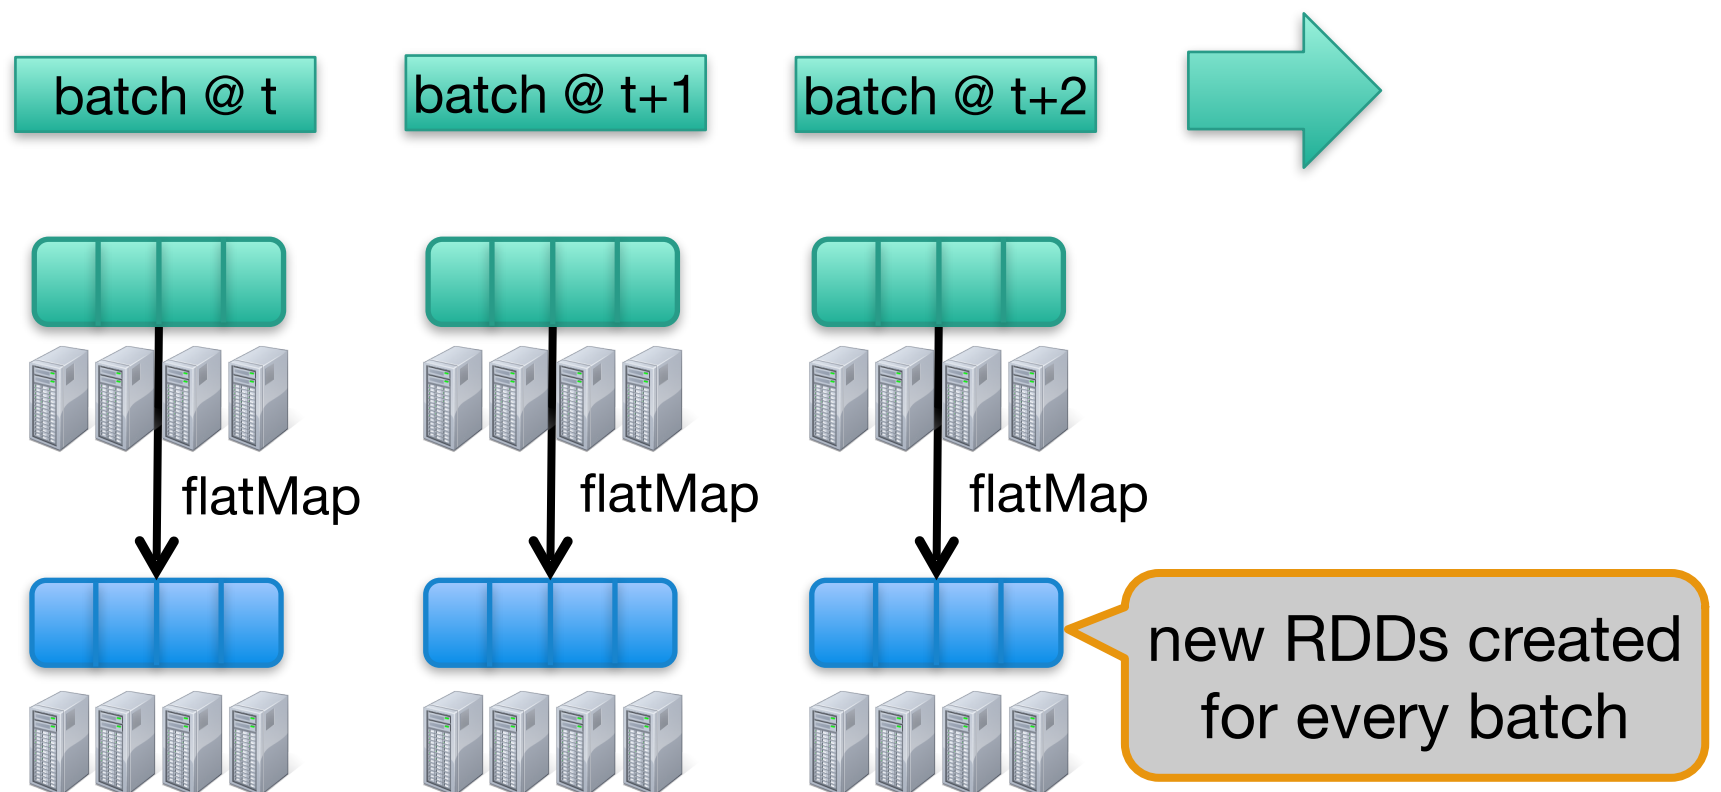
```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap(status => getTags(status))
```

new  
DStream

**transformation:** modify data in one DStream to  
create another DStream

tweets DStream

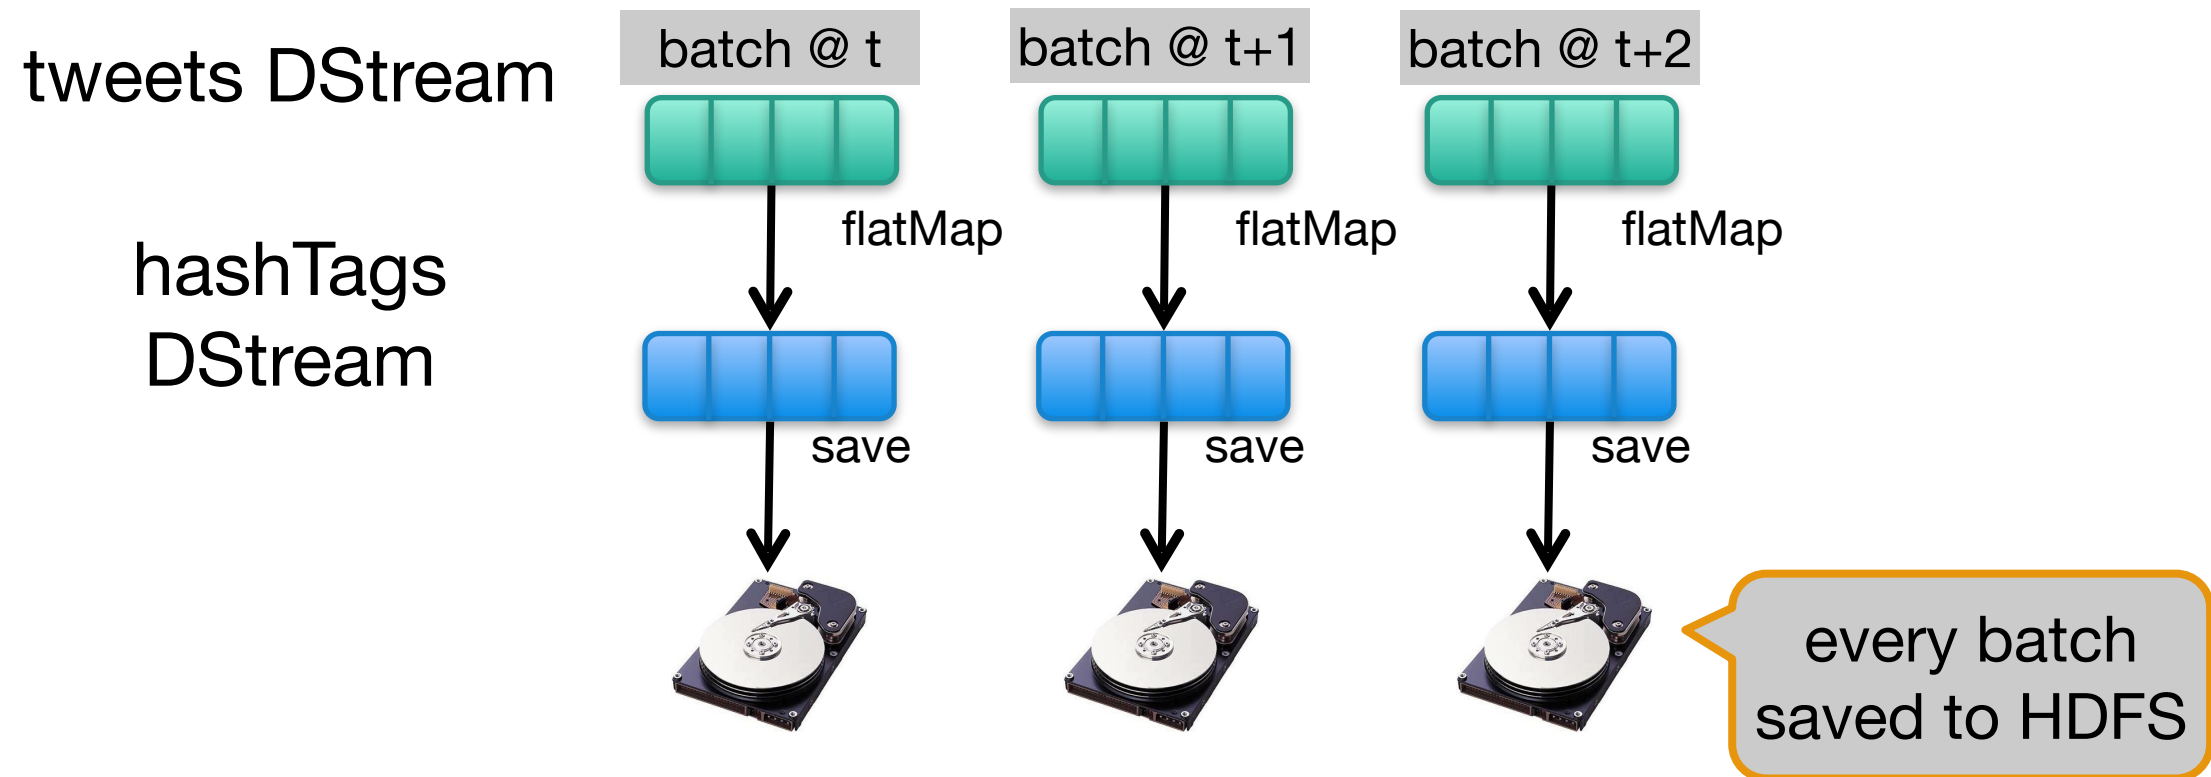
hashTags  
Dstream  
[#cat, #dog, ...]



# Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()
val hashTags = tweets.flatMap(status => getTags(status))
hashTags.saveAsHadoopFiles("hdfs://...")
```

**output operation:** to push data to external storage

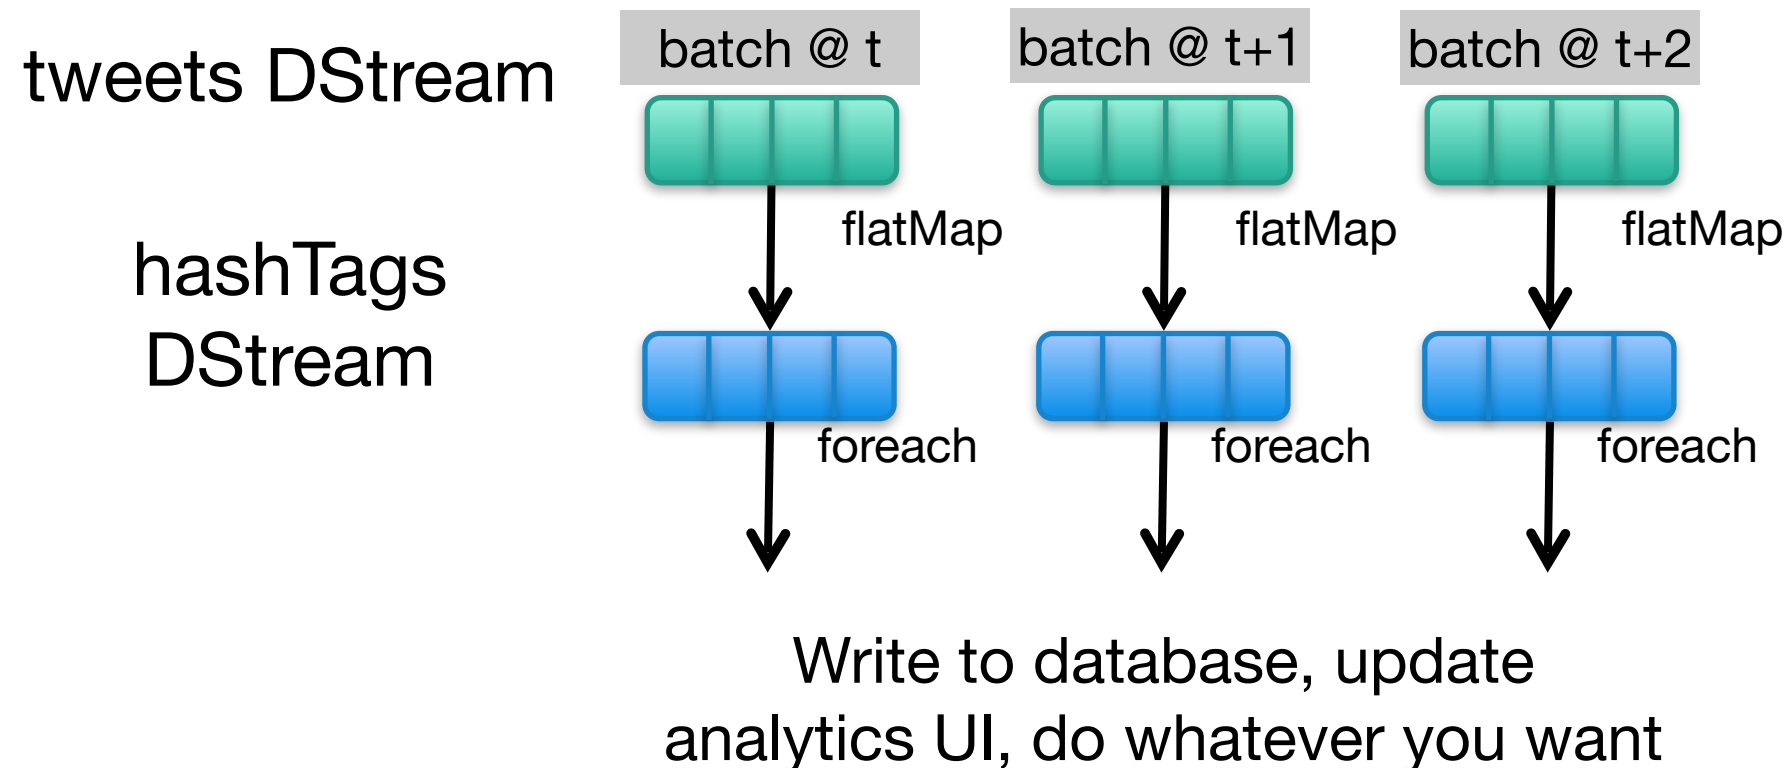




# Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()
val hashTags = tweets.flatMap(status => getTags(status))
hashTags.foreach(hashTagRDD => { ... })
```

**foreach:** do whatever you want with the processed data



# Java Example

---

## Scala

```
val tweets = ssc.twitterStream()
```

```
val hashTags = tweets.flatMap(status => getTags(status))
```

```
hashTags.saveAsHadoopFiles("hdfs://...")
```

## Java

```
JavaDStream<Status> tweets = ssc.twitterStream()
```

```
JavaDStream<String> hashTags = tweets.flatMap(new Function<...> { })
```

```
hashTags.saveAsHadoopFiles("hdfs://...")
```



Function object

# Window-based Transformations


```
val tweets = ssc.twitterStream()
val hashTags = tweets.flatMap (status => getTags(status))
val tagCounts = hashTags.window(Minutes(1), Seconds(5)).countByValue()
```

sliding window  
operation

window  
length

sliding  
interval

DStream of data



# Arbitrary Stateful Computations

---

Specify function to generate new state based on previous state and new data

- ▶ Example: Maintain per-user mood as state, and update it with their tweets

```
updateMood(newTweets, lastMood) => newMood  
moods = tweets.updateStateByKey(updateMood _)
```

# Arbitrary Combinations of Batch and Streaming Computations

---

Inter-mix RDD and DStream operations!

- ▶ Example: Join incoming tweets with a spam HDFS file to filter out bad tweets

```
tweets.transform(tweetsRDD => {  
tweetsRDD.join(spamHDFSFile).filter(...)  
})
```

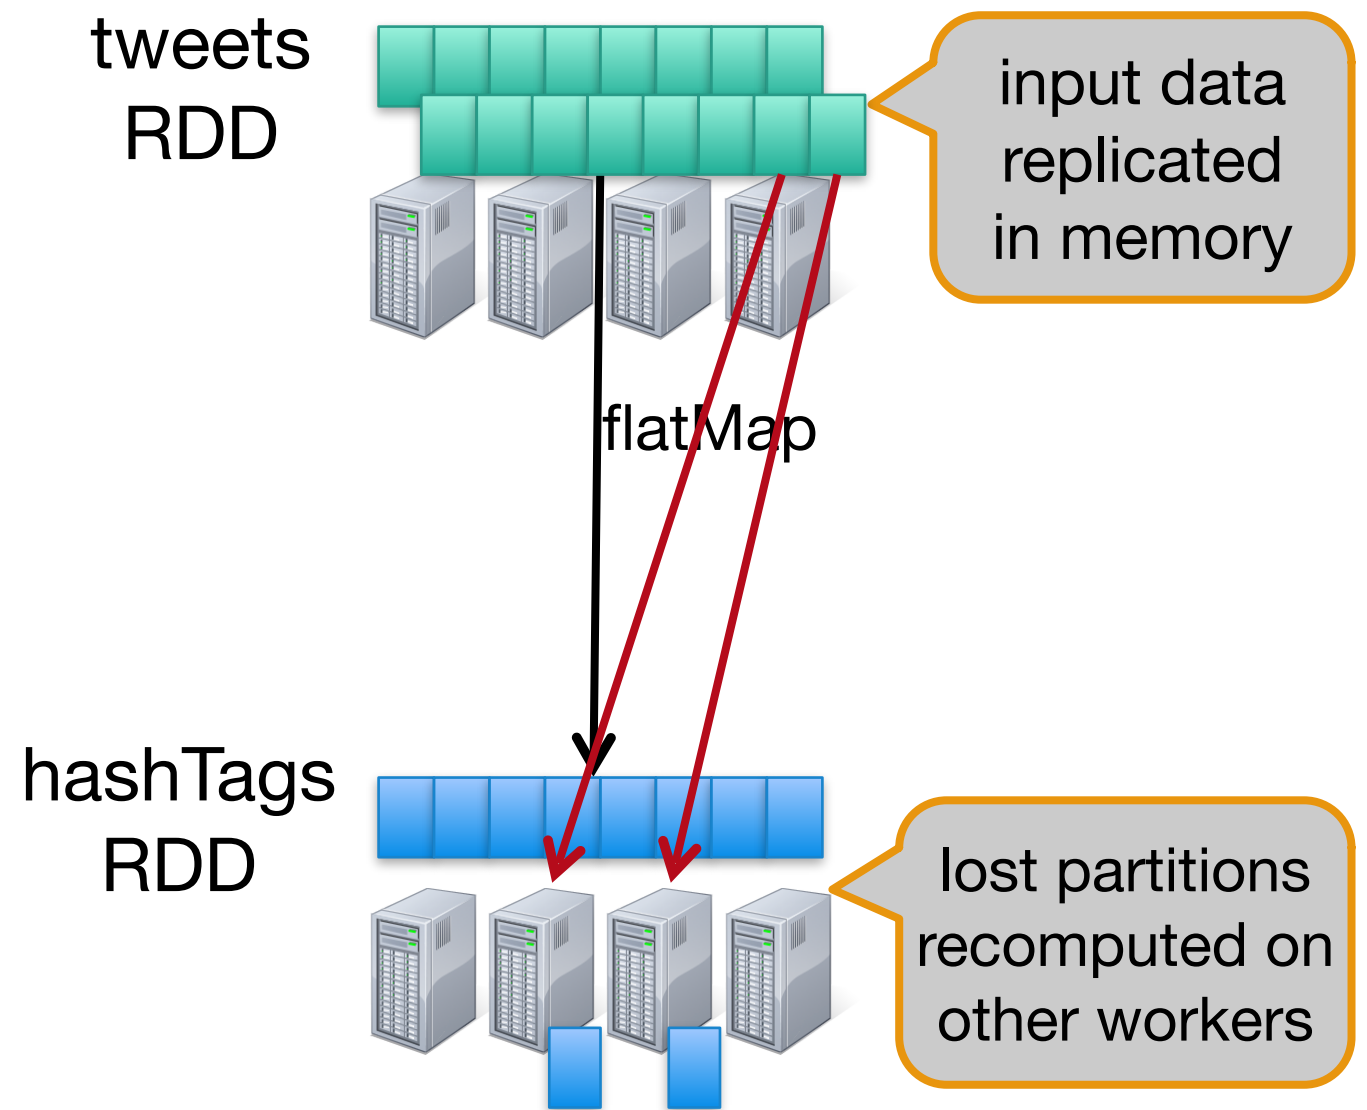
# DStream Input Sources

---

- ▶ Out of the box Spark Streaming provides
  - ▶ Kafka
  - ▶ HDFS
  - ▶ Flume
  - ▶ Akka Actors
  - ▶ Raw TCP sockets
- ▶ Very easy to write a *receiver* for your own data source

# Fault-tolerance: Worker

- ▶ RDDs remember the operations that created them
- ▶ Batches of input data are replicated in memory for fault-tolerance
- ▶ Data lost due to worker failure, can be recomputed from replicated input data



- All transformed data is fault-tolerant, and exactly-once transformations

# Fault-tolerance: Master

---



- ▶ Master saves the state of the DStreams to a checkpoint file
  - ▶ Checkpoint file saved to HDFS periodically
- ▶ If master fails, it can be restarted using the checkpoint file
- ▶ Automated master fault recovery coming soon



# Performance

Can process **6 GB/sec (60M records/sec)** of data on 100 nodes at **sub-second** latency



- ▶ Tested with 100 text streams on 100 EC2 instances with 4 cores each



# Comparison with Storm

Higher throughput than Storm

- ▶ Spark Streaming: **670k** records/second/node
- ▶ Storm: **115k** records/second/node



# Fast Fault Recovery

Recovers from faults/stragglers within **1 sec**

